

Agenda:

- What is a Database
- Why we are learning DB at Scaler.
- Type of Databases.
- Intro to Relational Databases
- Intro to Keys.

MySQL Databases.

Install MySQL Workbench + MySQL DB.

What is a Database. :-

Data :- Information.

Excel | Word | Notepad / . . .

Name	Phone no
Abc	1234
XY2	5678

⇒ Scalar's DB

→ Students

→ Instructors

→ Students.

Name	Email	batch	PSp	attendance
Nikhil	abc@—	123	97	95

batch: 123.

→ Name: Mar23 Intermediate

→ Start-date: 25th Mar.

→ Current Module: SDC.

====

⇒ Students, batches, Instructors, . . .

Students.csv

batches.csv

Instructors.csv

Companies.csv.

=====

Students.csv

Name	email	batch	PSP	Attendance	Coins	rank
Parth	abc	123	15	94	1000	5
Himanshu	def	432	96	92	2000	4
—	—	—	—	—	—	—

Ex: Avg PSP of batch = 123.

→ Iterate over the file row by row.

⇒ Issues of storing data in files.

- 1) Inefficient.
 - 2) Data Integrity.
- Student.txt:

Name	email	batch	PSP	Attendance	Coins	Phone
Parth	abc	123	15	94	1000	<u>abc.</u>
Himanshu	def	432	98	92	2000	—
—	—	—	—	—	—	—

CSV.

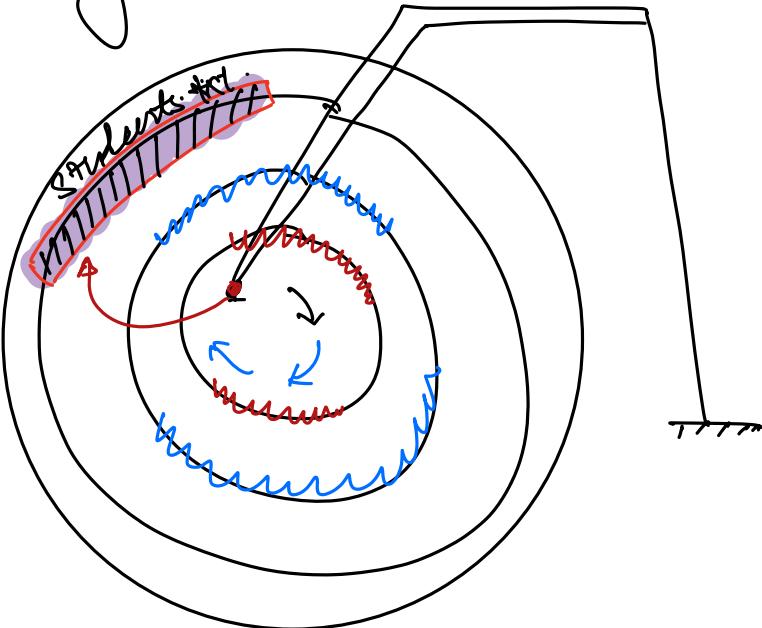
Deepak, Kulkarni, 9015608331, - - - -
 ↑ ↑ ↑
first Second Phone ... - - -

Rahul, Nandanwar, - - -
 ↑
Second.

? X

file → HDD.

3) Concurrency



Students.txt

(1)

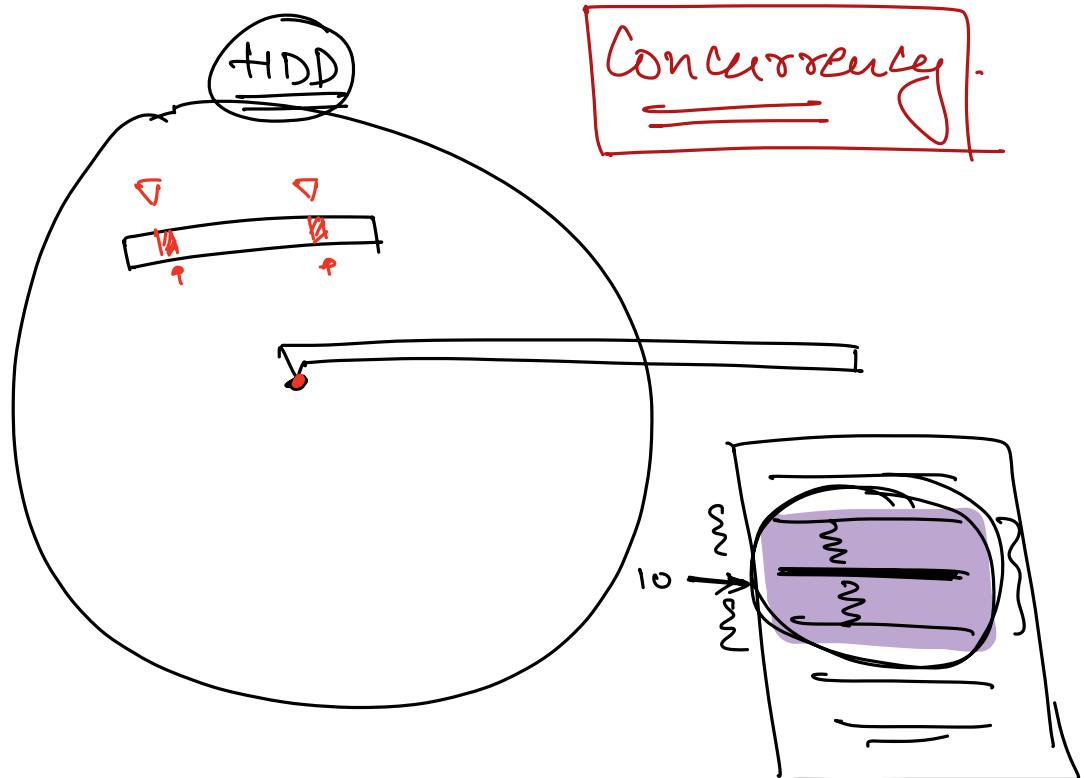
→ Only one user can read the data from HDD at one time.

⇒ Concurrency

⇒ Not fast enough.

file

10M



4) Security.

→ Role based access 

⇒ Database

Airbase : Place to store aircrafts.

Database : Place to store data.

⇒ DBMS : Database Management System

Type of Databases.

→ * Relational Databases.

→ Non Relational Databases.] NRD.

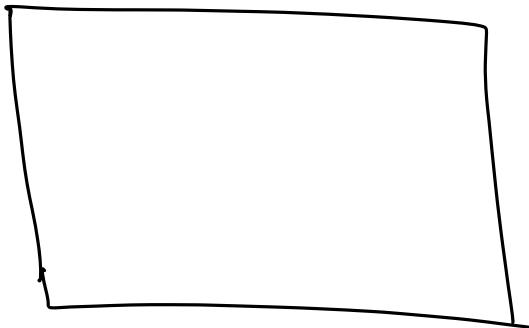
- Graph based (GraphQL).
- Mongo DB
- Columnar DB.
- Key-value DB.

Relational Databases.

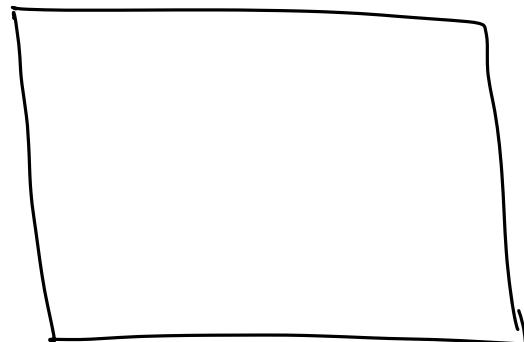
→ Collection of Tables
entity

⇒

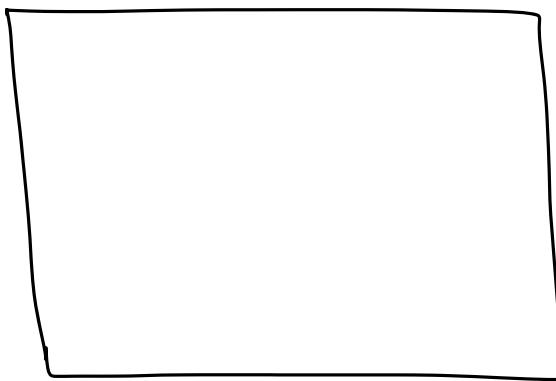
Students



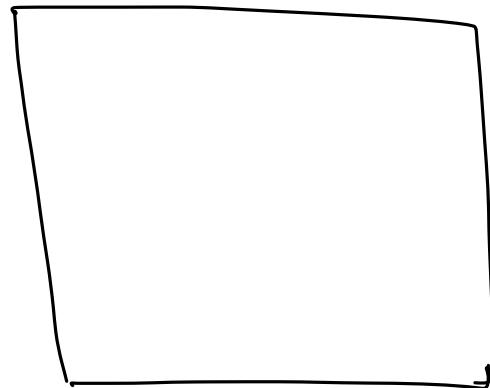
batches



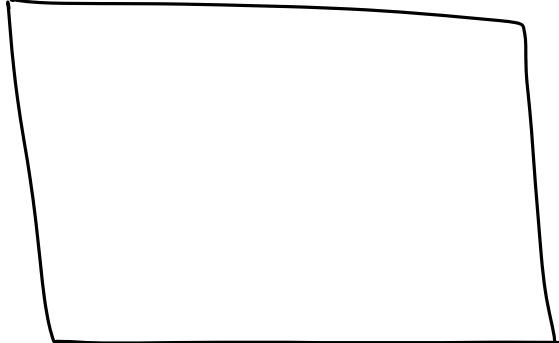
instructors



mentors



ta



...

⇒ Mathematics.

⇒



\Rightarrow Properties of Relational DB:

1) Every row is unique.

Students		
Rollno	Name	PSP
1	Deepak	95
2	Deepak	95
:		

2) All the values present in a column should be of same datatype.

3) Every cell should only contain atomic value.

\downarrow
Single/
Can't be divided
further.

\rightarrow Array, list, Map X

\Rightarrow Efficiency.

<u>4B</u>	<u>int.</u>
100	104
104	108
108	112

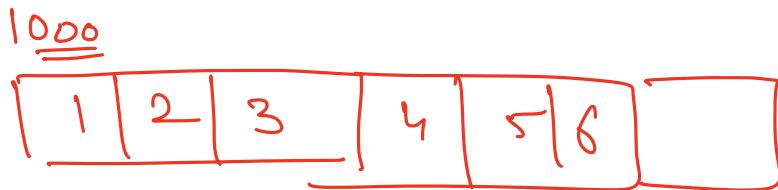
$O(1)$

$$100 + 4 * 3$$

112

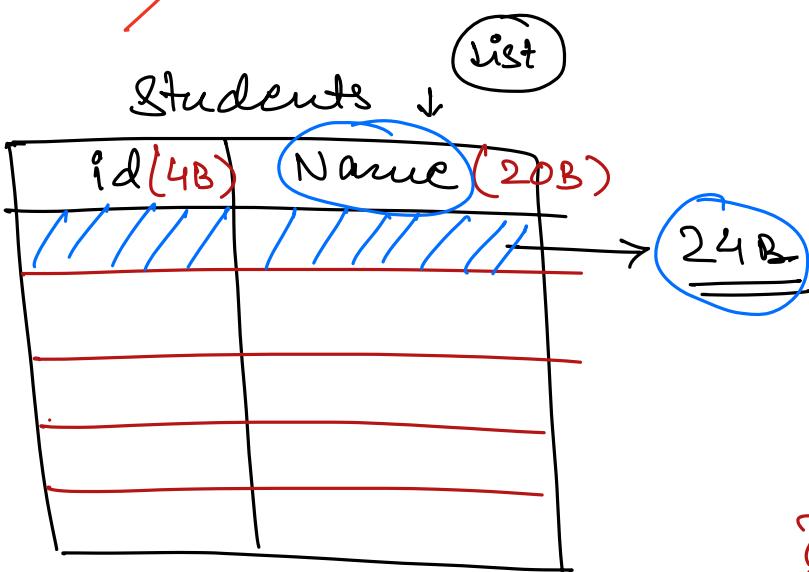
4B

1	2	3
4	5	6
7	8	9
10	11	12

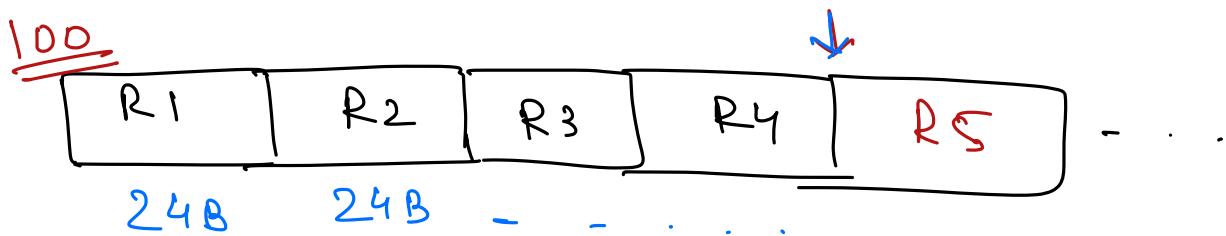


$$1000 + \underline{10 \times 4} \Rightarrow \underline{\underline{1040}}$$

...



?



$$100 + 24 \times 5 = \underline{\underline{220}}$$

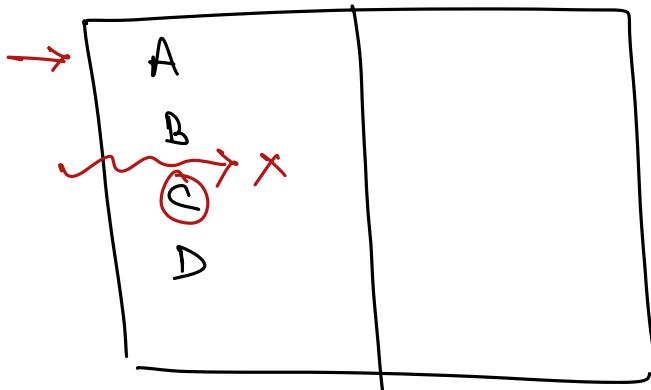
4) Order of columns is not guaranteed.

id	Name	age	batch	esp.
50	Ishan	25	1	95

⇒ Select * from Students
where id = 50.

⇒ Select Name, age, batch, esp from
Students where id = 50.

5) Order of rows is not guaranteed.

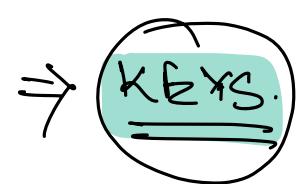


ORDER BY. ←

7) Column names should be UNIQUE.

Name	Name
Ramya P	Irin

Select Name from Students.



Students

Name	age	PSP
Venkatesh	25	95
Venkatesh	25	95

 : Column or a group of columns that can uniquely identify a row.

→ Super

→ Primary

→ Candidate

→ Composite

→ foreign.

SUPER KEY:

A column or a group of columns that can uniquely identify a row.

Students

name	PSP	email	batch	Phone-no

- 1) name X
- 2) batch X
- 3) email ✓
- 4) Phone no ✓
- 5) (email, phone) ✓
- 6) (name, email) ✓
- 7) (name, phone) ✓
- 8) (phone, PSP) ✓
- 9) (batch, PSP) X

Agenda:

⇒ keys

- Candidate Keys
- Primary Keys
- Composite Keys
- foreign keys.

⇒ Introduction to SQ L.

CANDIDATE KEYS.

Student

name	PSP	email	batch	Phone-no

- 1) name X
 - 2) batch X
 - 3) email ✓
 - 4) Phone no ✓
 - 5) (email, phone) ✓
 - 6) (name, email) ✓
 - 7) (name, phone) ✓
- }
- SK

- 8) (phone, PSP) ✓
 9) (batch, PSP) ✗

⇒ In many of the Super keys, there are columns that don't really play any role in uniquely identifying a row in a table. Even without these attributes we can uniquely identify a row.

Candidate key:

- ⇒ Super key of minimum size.
- ⇒ Super key where if we remove any of the attribute, then remaining part will not be a Super key.
- ⇒ Super key where all the attributes in that are required to uniquely identify a row.

(email, phone). → Sk ✓
 ↓
Ck ✗

(name, reg, batch-id) ⇒ Sk ✗
 Garima, 90, 123
 Garima, 90, 123

Ex

instructors

first-name	last-name	salary

Instructors can't have same name.

(first.name, last.name, salary) \Rightarrow SK ✓

CKX
=

(first.name, last.name) \Rightarrow SK ✓

CK ✓

\Rightarrow Ex : Attendance.

→

student-id	Class-id	attendance
1	12	80
2	12	90
3	12	70
4	14	85
5	90	95
6	91	95
7	20	85
8	20	85

Super Keys.

SK

- 1) Student-id \times
- 2) Class-id \times
- 3) (Student-id, Class-id) ✓
- 4) (Attendance) \times
- 5) (Student-id, Attendance) \times
- 6) (Class-id, Attendance) \times
- 7) (Student-id, Class-id, Attendance) ✓
- 8) (Student-id, Class-id, - - - -) ✓



Candidate Key
= ≠

(Student-id, Class-id)

⇒ Super Key of Min Size.



CK ⇒ Subset of SK

\Rightarrow EmpId ✓

Email ✓

(first-name, last-name) \Rightarrow SK X }

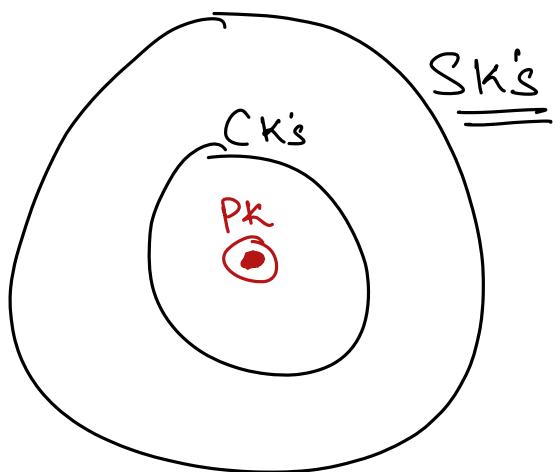
(last-name, department) \Rightarrow SK X }

\Rightarrow Super Keys \Rightarrow Party workers.

Candidate keys \Rightarrow MLA's.

Primary key \Rightarrow CM

Example-



PRIMARY KEY :

A Candidate key selected to be the unique set of columns to help identify a row in the DB.

Students

first-name	last-name	email	batch-id

Sk's \Rightarrow (email, first-name)
(email, batch-id)
(first-name, batch-id) \times

CK's. \Rightarrow email

↓
Pk.

\Rightarrow When we create a table in the DB,
then DB forces us to specify a column or
set of columns that will uniquely identify
a row in the table.

\Rightarrow Pk.

Properties of Pk.

1) Database Sorts the table on Pk.

Students.

St-id	name	email	Phone	PSP
1				
2				
3				
4				
5				
6				
7				

- 2) Database Outputs the data in the sorted Order based on the Pk.
- 3) Database creates index on Pk by default.

A good Pk:

- 1) Should be easy to sort on.
- 2) Should be small in size.
- 3) Should not change.

Students.

A hand-drawn diagram of a database table for students. The table has four columns: name, email, phone, and psp. The 'email' column is highlighted with a light blue box. A red circle at the top center contains the letters 'PK' underlined twice, with a red arrow pointing down to the 'email' column header. The data in the table is as follows:

name	email	Phone	psp
Kevin	Kevin@gmail.com	—	—
—	—	—	—
—	—	—	—

Can users change their email-id? Yes.

⇒ Changing the PK isn't good at all for your DB.

⇒ Ideally we shouldn't have user attributes as PK.

⇒ Always it is suggested to have a single integer (defined by the DB owner) as a PK.

Composite Keys.

⇒ Super Key with more than one attribute is a composite key

foreign key

Students.

st-id	name	email	Phone	batch-id
1	Parag	—	—	123
2	Venkatesh	—	—	123
3	Soham	—	—	123
4	Deepran	—	—	456

FK.

batches.

batch-id	name	start-date	instructor	no. of students.
123	Morning	—	—	100

A column or a group of columns that helps us to uniquely identify a row in the other table.

⇒ Students

st-id	name	batch_id
1	ABC	100
2	XYZ	10
3	PQR	11
4	DEF	11

batches

batch_id	batch_name
100	Mar23 Sat
11	Mar23 Mon
12	Mar23 Beg
13	
14	

① If we are adding a student to a batch, the it should exist.

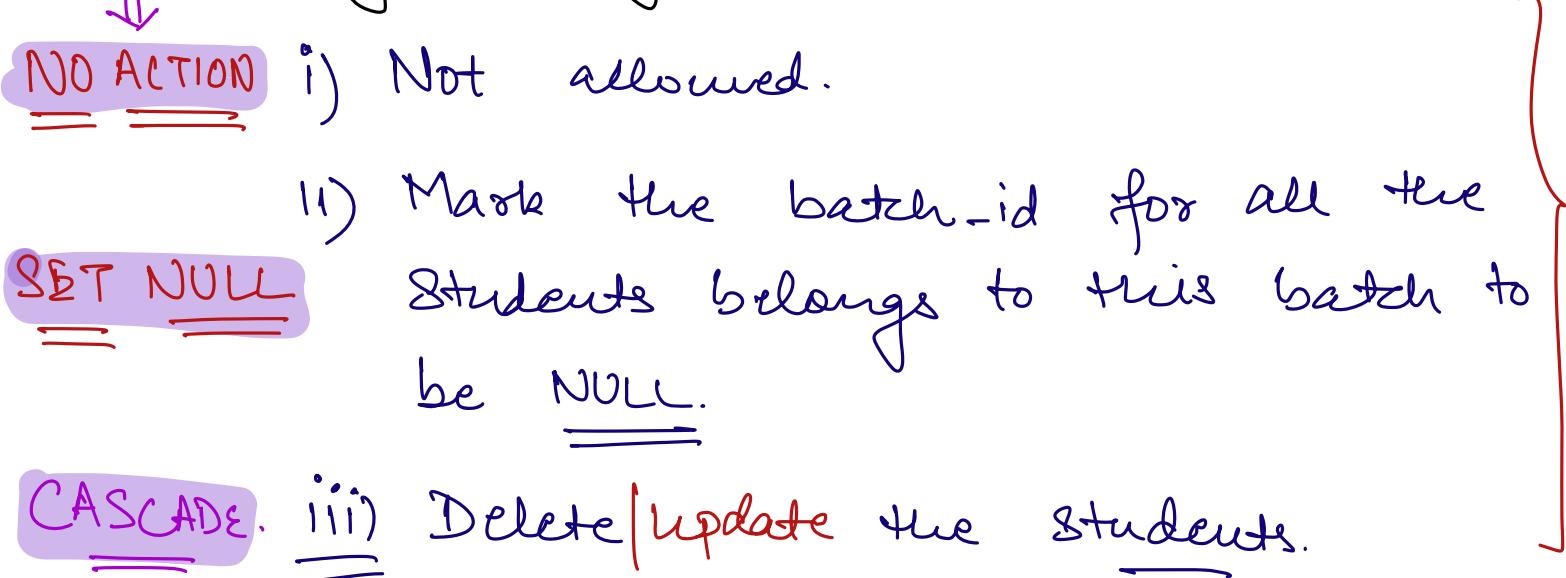
② Deleting/Updating a batch.



NO ACTION i) Not allowed.

ii) Mark the batch_id for all the students belongs to this batch to be NULL.

CASCADE. iii) Delete/Update the students.



→ book marks.

id	name	class-id
1	—	10
2	—	10
3	—	10
:	—	10

id	desc
x 10	—

Ex

Student

St-id	name	b-id
1	Rahul	10
—	—	10

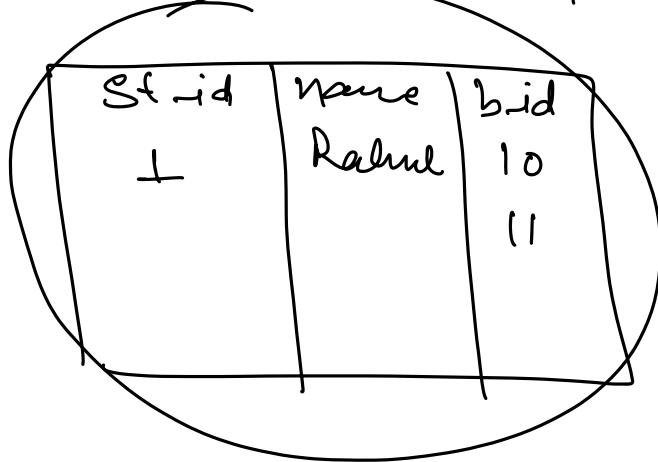
batches ↓

b-id	b-name
10	Mar23 Intermediate
11	Feb.
10	Mar23
10	—

Ex

Student

⇒ 10



b-id	b-name
10	Mar23 Intermediate
11	—
12	—
13	—

(NO ACTION)

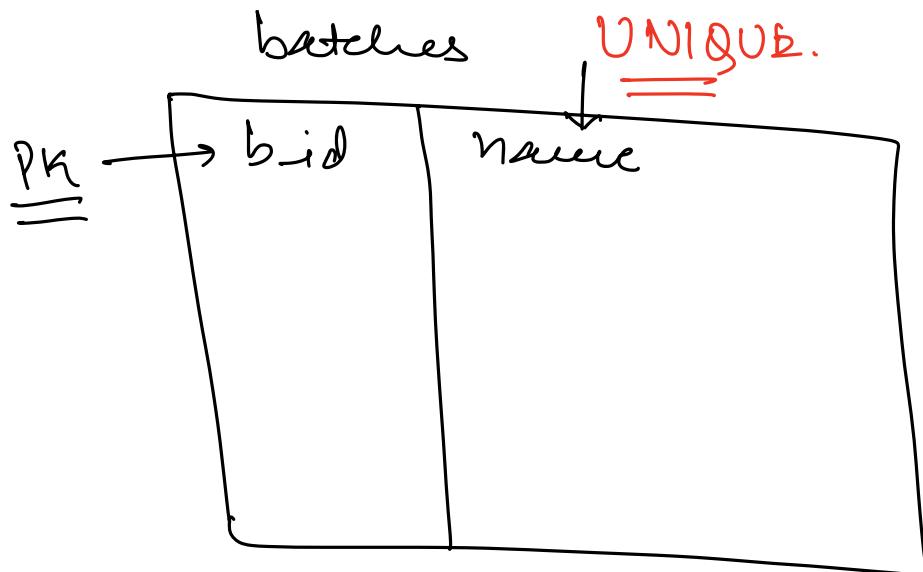
Delete the batch with batch-id. → Not Allowed

(SET NULL)

Mark b_id as
NULL in students.

→ Do what
we did in
the other table
(CASCADE)

we did in
the other table
(CASCADE)



\Rightarrow (St-id, f-name, l-name) SK ✓
Composite Key ✓

```
SELECT * FROM film;
-- Print
-- 2023-10-16 08:06:20
INSERT INTO film
VALUES
(default, "Tiger", default, 2012, 1, NULL, 3, 15, 120, 10, "G", "Trailers,Behind the Scenes",
default);

-- Error Code: 1048. Column 'last_update' cannot be null
-- Error Code: 1364. Field 'title' doesn't have a default value
```

```
select film_id as id, title as film_name
from film
order by film_name;
```

```
-- DISTINCT
```

```
select distinct rating from film;
```

```
select distinct rating, release_year
from film;
```

```
select "Hello World!";
```

```
select title, "Hello World" as greeting
from film;
```

```
-- Print the title and duration of the movie (in HRS)
```

```
select title, round(length/60) as length_in_hrs
from film;
```

```
create table film_copy (
    title VARCHAR(128) NOT NULL,
    release_year YEAR DEFAULT NULL
);
```

```
select * from film_copy;
```

```
insert into film_copy
select title, release_year from film;
```

```
select distinct description  
from film;
```

```
select distinct title  
from film;
```

```
select * from film  
where release_year = 2006 AND rental_duration < 5;
```

-- NOT

-- Get all the movies which were not released in 2006

```
select * from film  
where release_year <> 2006;
```

-- Get all the movies with rating "G" OR "PG"

```
select * from film  
where rating = "G" or rating = "PG";
```

-- IN Operator.

```
select * from film  
where rating NOT IN ("G", "PG");
```

Agenda :

READ .

- BETWEEN
- LIKE
- IS NULL
- ORDER BY
- LIMIT.

UPDATE

DELETE .

⇒ BETWEEN.

Find all the students with PSP ≥ 50 & $PSP \leq 70$.

```
Select *
from Students
where PSP >= 50 AND PSP <= 70.
```

Find all the movies released b/w 2010 & 2020.

```
Select *
from film
where release-year >= 2010 and
release-year <= 2020.
```

⇒ BETWEEN a AND b ⇒ $>= a$ AND $\leq b$.

⇒ Select *
from Students
where PSP between 50 AND 70;

⇒ Select *
from film
where release-year BETWEEN 2010 AND 2020;

⇒ BETWEEN 20 AND 50

⇒ 20	✓
⇒ 19	✗
⇒ 31	✓
⇒ 50	✓
⇒ 51	✗

⇒ BUILT-IN functions.

- Recordings.
- Typed Notes.
- Google.

⇒

⇒ LIKE.

↳ String Pattern matching.

Students.

id	name	batch_name				
1	Venkat	Academy	Mar22	1st	Morning	
2	Priya	Apr22	beg		Academy	Morning

⇒ All the academy batches contains Academy in the name.

⇒ Morning batches contains Morning in the name.

Q. Return all the morning batches in Academy.

```
bool checkMorningAcadBatches(String bName){  
    if(bName.contains("Morning") &&  
        bName.contains("Academy")) {  
        return true;  
    }  
}
```

3

return false

3
=

$\Rightarrow \underline{\text{LIKE}} : \underline{\text{STRNG}} \quad \underline{\text{MATCHING}}$.

$\Leftrightarrow \underline{\text{wildcard}} \quad \underline{\text{characters}}$.

'%': Any character any # of times.
(0, +, 10, 20, ...)

'_': Any character exactly + time.

\Rightarrow Select *
from film
where desc LIKE '%Hello.%'

\Rightarrow Hello ✓
Hello
Hello

$\Rightarrow \underline{\text{Cat}} : \underline{\text{--t}} \quad \checkmark$
%t $\quad \checkmark$
. $\quad \checkmark$
%. Cat %. $\quad \checkmark$

Q. Return all the morning batches

Select *

from batches

where batch-name LIKE '%Morning%';

Q. Return all the morning batches in Academy.

Select *

from batches

where batch-name LIKE '%Morning%'
AND LIKE '%.Academy%';

⇒ — Morning — Academy ✓

— Academy — Morning — ✓

⇒ LIKE '%.academy%'

ACADEMY Academy
✓ ✓

⇒ Students with mes in their names.

Select *

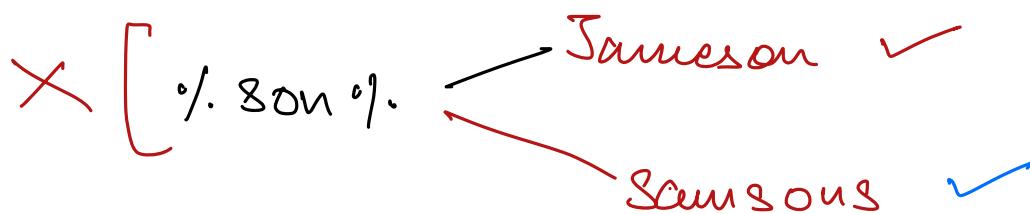
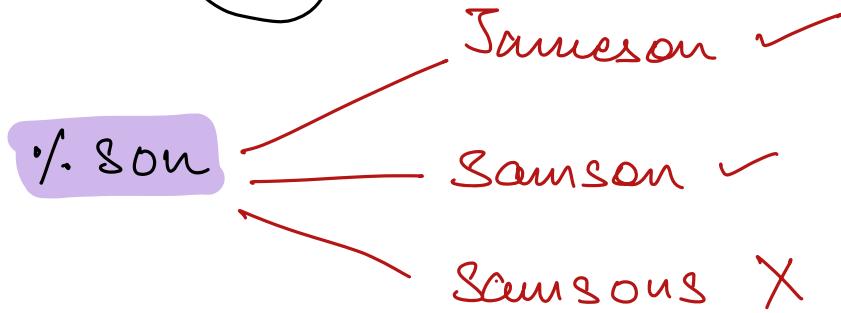
from students

where name LIKE '%.mes%'

\Rightarrow Animesh ✓

\Rightarrow Umes ✓

\Rightarrow Ending with son.

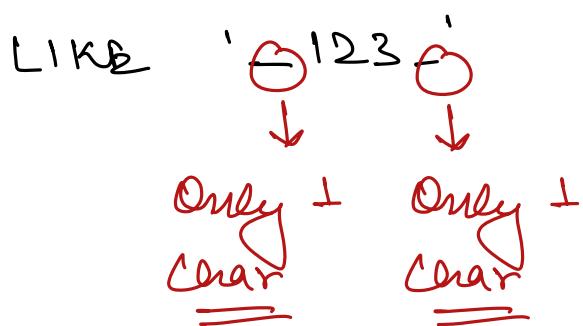


\Rightarrow Contains the word 'moon'

%moon%

\Rightarrow

 1 2 3

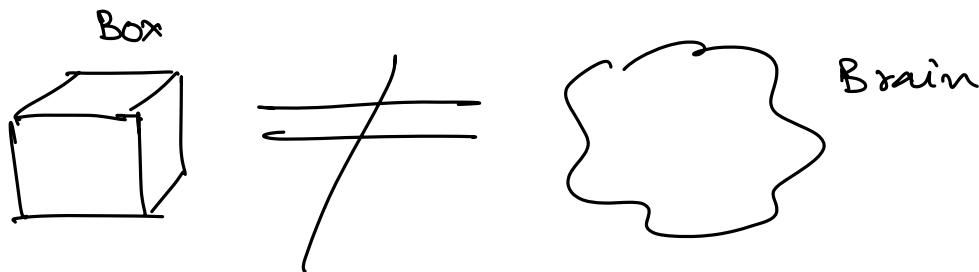


~~%123%~~ abc123xyz ✗

IS NULL

⇒ We can't use = operator to match NULL
Values.

⇒ NULL = NULL X



⇒ IS NULL.

Select *
from film
where title = NULL X

Select *
from film
where title IS NULL;

ORDER BY

⇒ By default the order of data is NOT
guaranteed.

⇒ If want the final ans to be ordered by
some column

⇒ ORDER BY

Select *

from film

⇒ ASCENDING.

Order by title;

Select *

from film

⇒ DESCENDING.

Order by title Dec;

Select *

from film

Order by releaseyear, name;

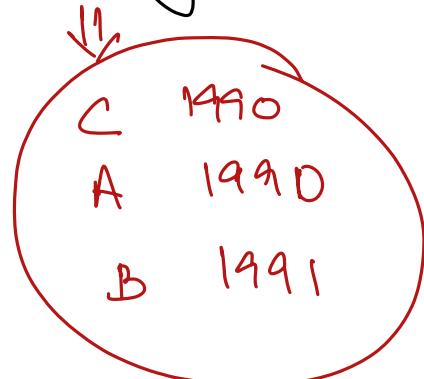
Students.

id	name	yob
1	C	1990
2	B	1991
3	A	1990

Select *

from students

Order by yob



Select *

from students

Order by yob, name;

A 1990 }
C 1990 }
B 1991 }

NOTE: Ultimate tie breaker is always PKs

table : [{ } { } - - - { }]

ans : []

for each row in table :

if condition is matched :

ans.add(row);

ans.sort (sorting condⁿ);

ORDER BY with DISTINCT keyword

Select DISTINCT title

from film

ORDER BY release-year ;

⇒ Students.

Ex

id	name	YOB
1	abhishek	1998
2	abhishek	1990
3	bhim	1995

Select distinct name
from Students
Order by job;

⇒ Bhim. 1995
Abhishek 1998

If we have distinct in select clause, then
we can only order by the columns present
the distinct clause.

Select distinct a, b

from table

Order by a



Select distinct a, b

from table

Order by c



distinct
Select *
from —

where name = —

Order by —

FROM
↓
WHERE
↓
ORDER BY
↓
DISTINCT
↓
PRINT

=>

table : [{ } { } - - - . { }]

ans : []

for each row in table :

if condition is matched : \Rightarrow where

ans.add(row);

ans.sort(sorting cond")

final-ans = set(ans) \Rightarrow DISTINCT.

print(final-ans)

Ambiguity.

id	name	yob
1	Abhishek	1998
2	Abhishek	1990
3	bhim	1995

Select distinct name
from Students
Order by yob;

ans : [{ Abhishek, 1998 } { Abhishek, 1990 } { Bhim, 1995 }]



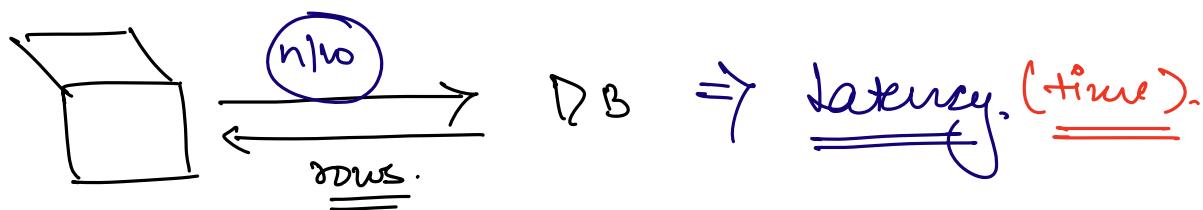
ans : [{ Abhishek, 1990 } { Bhim, 1995 } { Abhishek, 1998 }]

↓
DISTINCT NAMES.

\Rightarrow Set doesn't preserve the order of elements.

\Rightarrow LIMIT.

Select *
from Students. \Rightarrow 10⁶ rows.



\Rightarrow LIMIT.

Select *
from Students.
LIMIT 10;

}

first 10 rows

\Rightarrow 11 to 20.

Select *
from Students.
LIMIT 10 OFFSET 10;

C R U D

\Rightarrow Update.

\hookrightarrow Updating row data.

Update {table-name}

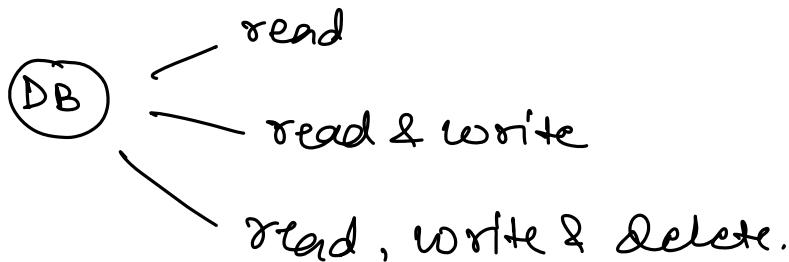
Set PSP = 99;

\Rightarrow Update the PSP
for all the
Students.

Update {table-name}

Set PSP = 99

Where id = 10.



DELETE.

\hookrightarrow Rows

delete from {table}

\Rightarrow Deletes all rows.

delete from {table}

Where id = 10



table : [{ } { } - - - { }]

for each row in table :

if row matches condⁿ:

 deletes row;

Agenda :

→ **JOINS.**

→ **Self Join.**

JOINS.

⇒ Till now only single table was involved in the queries.

⇒

students.

id	name	batchid
1	Vishal	1
2	Rahul	1
3	Deepak	2
4	Ajay	3

batches	
id	name
1	A
2	B
3	C
:	
10	

Q. for every student print their batch name.

name	b-name
Vishal	A
Rahul	A
Deepak	B
Ajay	C

⇒ We need to combine data from 2 tables.

↓
JOINS.

⇒ JOINS.: Allows us to combine the ^{data} across multiple tables.

SYNTAX.

Select Students.name, batches.name
from Students
join batches
ON student.batchid = batches.id;

Condition



Select s.name, b.name
from Students as s
join batches b
ON student.batchid = batches.id;



JOIN

Condition.

\Rightarrow ans: []

$O(NM)$

```

for each row in the Students:
    for each row in the batches:
        if join cond matches:
            ans.add(row);
print(ans);

```

JOINS: We are combining the data from Students & Batches table based on the batch_id.

\Rightarrow STUDENT BUDDY.

We assign a senior Student as a Buddy to each scalar Student.

Students

id	name	batch_id	Rsp	buddy_id
1	Parth	1	85	4
2	Umesh	2	90	4
3	Venkata	3	92	4
4	Pesarak	2	93	NULL

Buddy 17: self referential Column.

- ⇒ A buddy is nothing but a student itself.
⇒ Buddy id is an id of another student.

Q: for every student, print their name along with their buddy name.

Students

id	name	batchid	Rsp	buddy-id
1	Parth	1	85	4
2	Vinesh	2	90	4
3	Venkata	3	92	4
4	Pesarak	2	93	NULL

buddies

id	name

Select Students.name, buddies.name
from students
join buddies
ON Students.buddy_id = buddies.buddy_id;



Select S1.name as Name, S2.name as Buddy_name
from students S1
join students S2
ON S1.buddy_id = S2.id

⇒ SELF JOIN.

⇒ In self join, table name aliases are mandatory.

b.x. Employees.

employees.

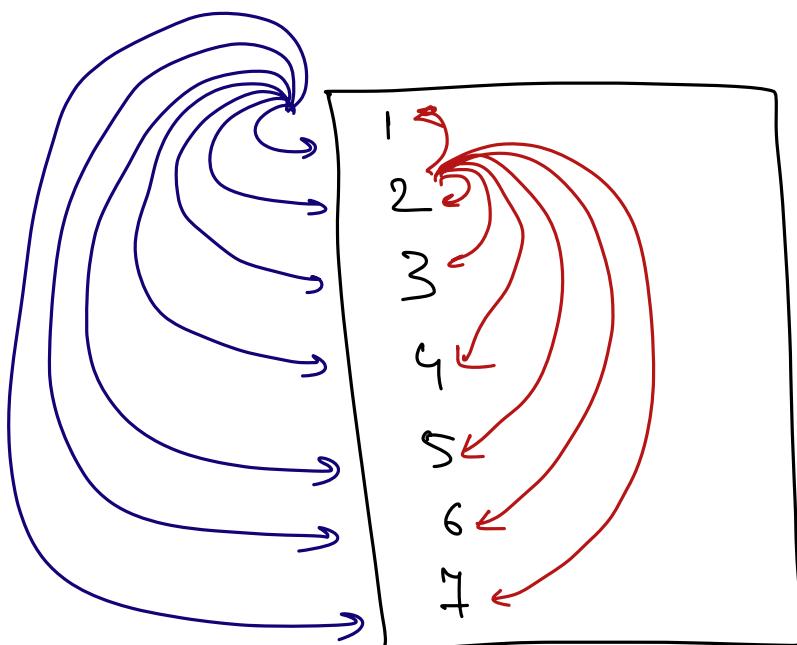
id	name	joining-date	department	manager-id

Q. for every employee, print its name & its manager's name.

Select e.name as Name, m.name as Manager
from employees e
join employees m
ON e.manager-id = m.id

ans: []

{ for each row in the Students:
 for each row in the Students:
 if join condⁿ matches:
 ans.add(row);
 print(ans);



JOINS.

Students

id	name	b-id
1	Rajesh	1
2	Umesh	1
3	Jasim	2

batches

id	b-name
1	A
2	B

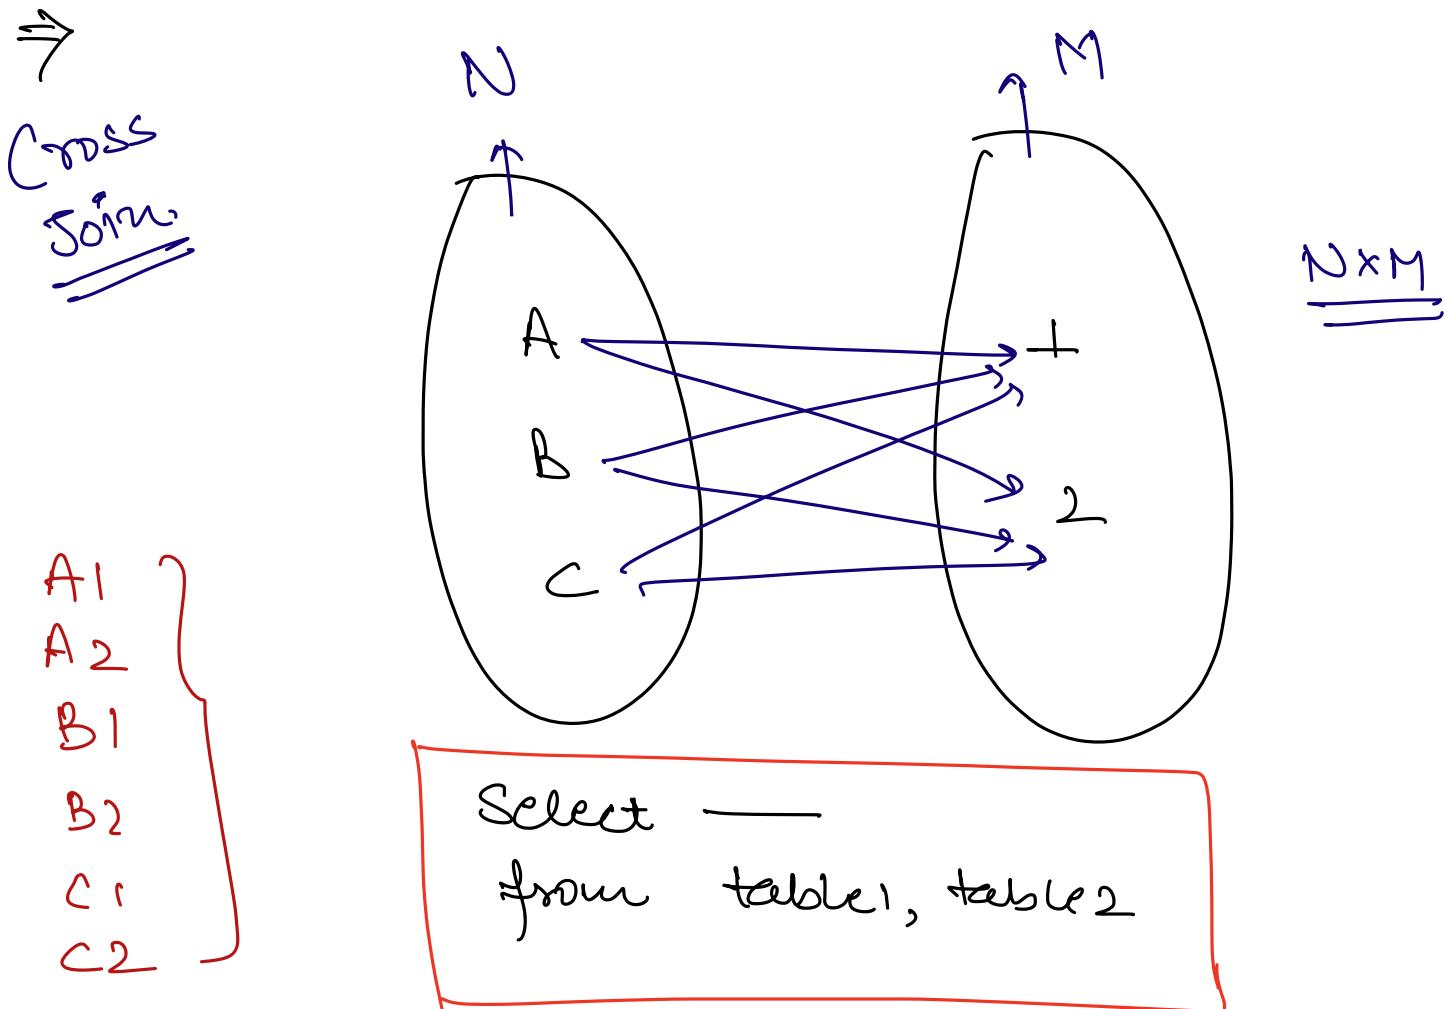
JOIN
Condⁿ

Students (S)			batches (B)	
id	s-name	b-id	id	b-name
1	Rajesh	1	1	A
2	Vinesh	1	1	A
3	Jasim	2	2	B

Temp Table.

Select s-name, b-name.

```
from Students S
join batches B
ON S.b-id = B.id
→ temp table.
```



Students

id	name
1	A
2	B
3	C

job_referrals.

id	referral
1	Amazon
2	Google.

St-name	referral
A	Amazon
A	Google
B	Amazon
B	Google
C	Amazon
C	Google

Select _____
from Students, referrals;

- Left Join
- Right Join
- Inner Join
- Outer Join.

Select s.name, b.name

from students as s

join batches b

ON s.batchid < b.id;

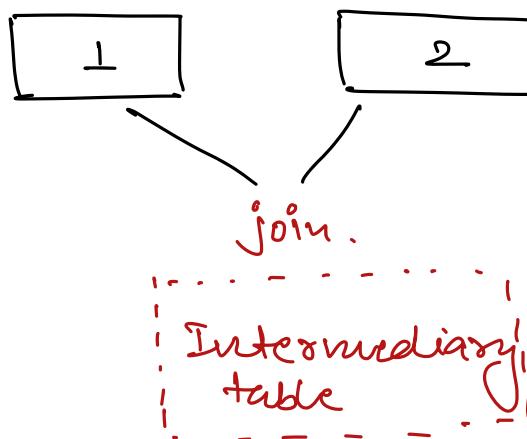


id	name	batchid
1	Vishal	1
2	Rahul	1
3	Deepak	2
4	Ajay	3

id	name
1	A
2	B
3	C

— * —

Joining multiple tables



Ex:

Students.

id	name	instructor_id	batch_id

instructors

id	name

batches

id	name

- Q. for every student, give their name along with the name of their instructor & name of their batch.

$$1 + 2 + 3 + 4 + 5 + \dots \dots$$

(1+2) + (3+4) + \dots \dots

\Rightarrow Join the tables in pairs.

Select s.name, i.name, b.name

from students s

join instructors i

ON s.instructor_id = i.id

join batches b

ON s.batch_id = b.id

Students (s)				instructors (i)		batches (b)	
id	name	instructor_id	batch_id	id	name	id	name

table1 : []
table2 : []
table3 : []
ans : []

ans1 : [] *// Intermediary-*

for each row1 in table1:

 for each row2 in table2:

 if condⁿ is TRUE:

 ans1.add (row1 & row2)

for each row in ans1:

 for each row in table3:

 if condⁿ is TRUE:

 ans.add (_____);

friend (ans);

COMPOUND JOINS.

Q. for every film, name all the films that were released with ± 2 years of that film & their rental rate $>$ rental rate of the movie.

film

Name	release year	rental rate
Sholay	2000	2
DDLJ	1999	3
KKHH	1998	2
3 Idiots	1994	4

⇒ Output

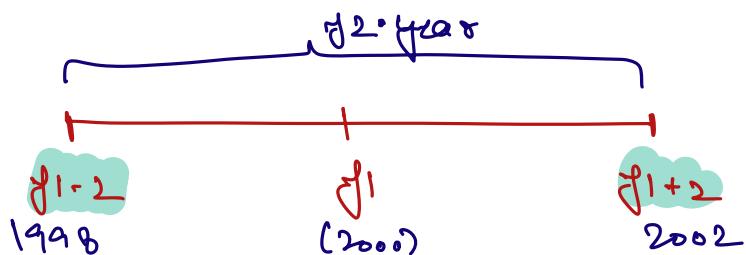
Sholay | DDLJ.
DDLJ | 3 3 Idiots.
KKHH | DDLJ
KKHH | 3 3 Idiots.

film (f_1)

Name	release year	rental rate
Sholay	2000	2
DDLJ	1999	3
KKHH	1998	2
3 Idiots	1994	4

film (f_2)

Name	release year	rental rate
Sholay	2000	2
DDLJ	1999	3
KKHH	1998	2
3 Idiots	1994	4



Select f1.name, f2.name
 from film f1
 join film f2
 ON (f2.year >= f1.year - 2 AND
 f2.year <= f1.year + 2) AND
 f2.rental_rate > f1.rental_rate;
 } BETWEEN.
 f2.year BETWEEN (f1.year - 2, f1.year + 2)

- ① Join need not to have always equality cond.
 - ② Join can also have multiple conditions in
ON clause. \Rightarrow Compound Join.
- \Rightarrow Type of Joins.

Students

id	name	buddy_id

Q. for every student,
 print their name
 along with their buddy
 name.

Students (S) $\xrightarrow{\text{Left}}$

id	name	buddy_id
1	A	2
2	B	NULL
3	C	2
4	D	1

Students (b) $\xrightarrow{\text{Right}}$

id	name	buddy_id
1	A	2
2	B	NULL
3	C	2
4	D	1

Select _____

from Students S

join Students b

ON S.buddy_id = b.id;

inner

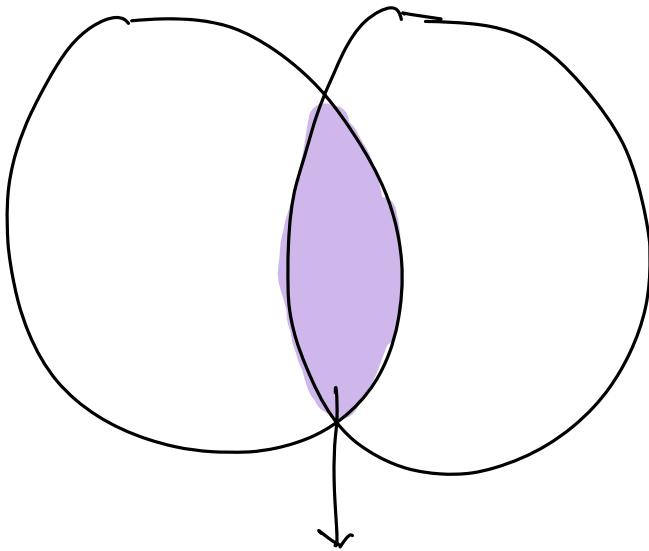
Optional.

Student Buddy

A | B

C | B

D | A



Inner
Join.

⇒ When do a join b/w 2 tables Left & Right,
 If a row of left table doesn't match with
 a row of right table, then we'll not have
 that row in the output & Vice-versa.

Q. for every student print their names along
 with names of their batches.

Students.

id	name	batch-id
1	John	1
2	Jane	1
3	Jim	Null
4	Jenny	Null
5	Jack	2

batches

id	name
1	A
2	B
3	C

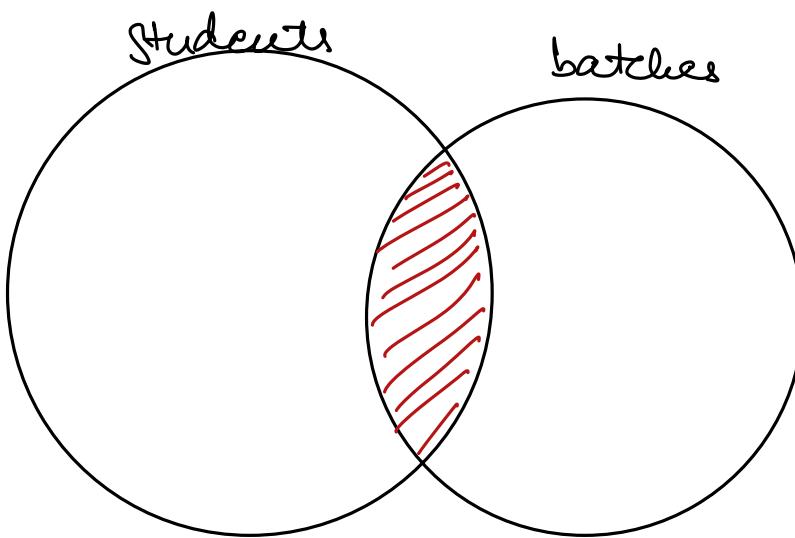
⇒

John | A
 Jane | A
 Jack | B

Select S.name, b.name
 from Students s
 join batches b
 ON S.batch-id = b.id ;

⇒ INNER
JOIN.

INNER JOIN : Join that includes only rows from both the tables that matches the condⁿ.



OUTER JOINS

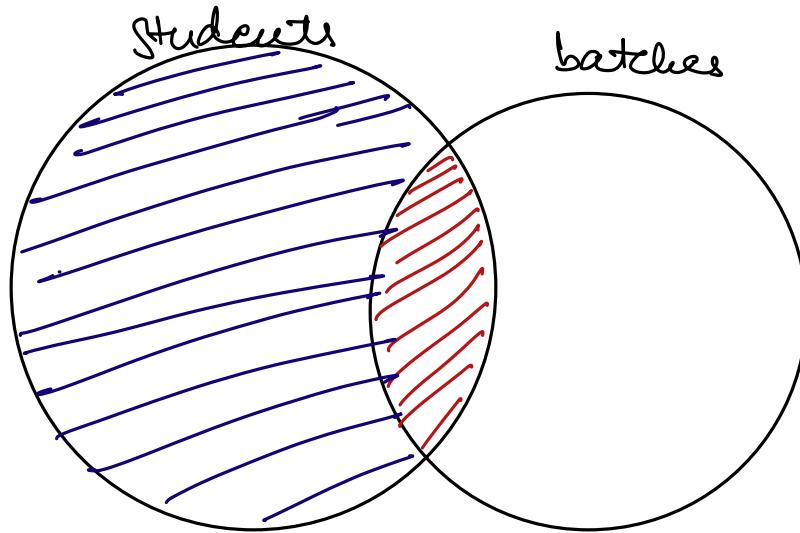
⇒ Allows us to get the rows which don't matches the Condⁿ.

- 1) Left JOIN. / Left Outer Join.
- 2) Right Join / Right Outer Join.
- 3) full Outer Join.

↳ MySQL doesn't support.

LEFT JOIN.

→



→ Includes all the rows matches the condⁿ
+

Includes all the rows from the left table
that doesn't match the condⁿ, for these rows
we'll put NULL for all the columns of
right table.

Students.

id	name	batch-id
1	John	1
2	Jane	1
3	Jim	NULL
4	Jenny	NULL
5	Jack	2

batches

id	name
1	A
2	B
3	C

John		A
Jane		A
Jack		B
Tim		NULL
Jenny		NULL

Select s.name, b.name
from students s

left join batches b
ON s.batch_id = b.id ;

RIGHT JOIN.

Includes all the rows matches the cond^

+

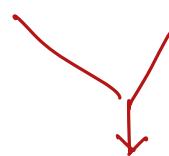
Includes all the rows from the right table
that doesn't match the cond^, for these rows
we'll put NULL for all the columns of
left table.

Students

id	name	batch.id
1	John	1
2	Jane	1
3	Jim	Null
4	Jenny	Null
5	Jack	2

batches

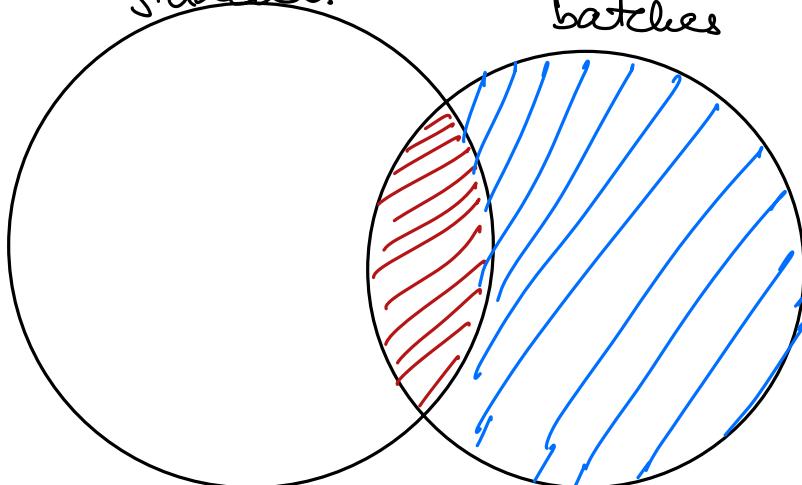
id	name
1	A
2	B
3	C



John		A
Jane		A
Jack		B
Null		C

Students

batches



FULL OUTER JOIN.

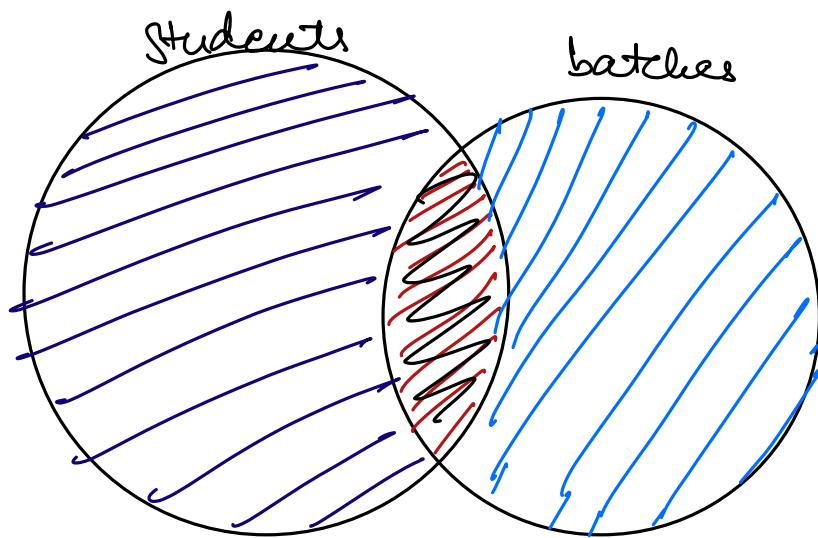
Includes all the rows matches the condⁿ
+

Includes all the rows from the left table
that doesn't match the condⁿ, for these rows
we'll put NULL for all the columns of
right table.

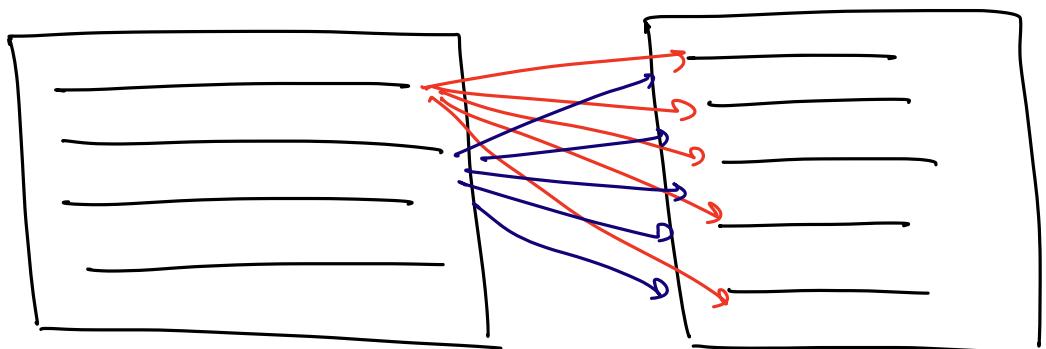
+

Includes all the rows from the right table
that doesn't match the condⁿ, for these rows
we'll put NULL for all the columns of
left table.

John		A
Jane		A
Jack		B
Jim		NULL
Jenny		NULL
Null		C



Cross Join.



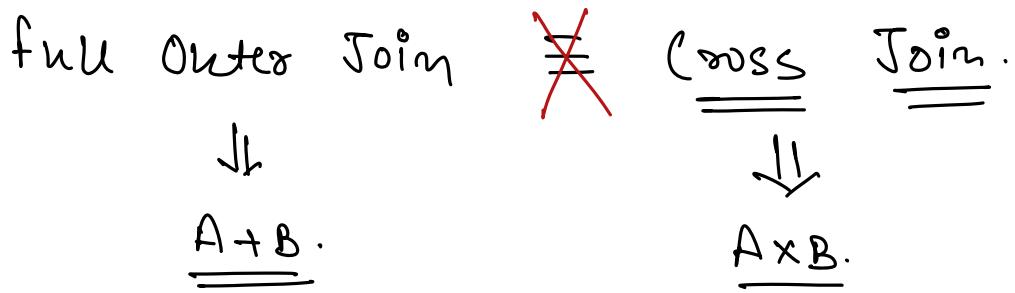
Select —
 from table1 \Rightarrow 9
 join table2 \Rightarrow 4
 ON true } \Rightarrow 36.
 9×4

$\Rightarrow O(N \times M)$

\Rightarrow No condn.

Select _____

from table1, table2;



USING.

- joining based on the equality condⁿ.
- Name of the column is same.

Select _____

from student s

join batches b

ON s.batch_id = b.batch_id



USING (batch_id)

NATURAL JOIN.

⇒ When we join tables based on equality condⁿ for all the same columns, that is Natural join.

table ①

a	b	c	d	e
---	---	---	---	---

table ②

b	c	x	r	e
---	---	---	---	---

Select _____

from table1

join table2

ON table1.b = table2.b AND

table1.c = table2.c AND

table1.e = table2.e

⇒ Select _____

from table1

join table2

USING (b, c, e);

Select _____

from table1

Join table2



Only joins based on the
same columns.

IMPLICIT JOIN \equiv Cross Join.

Select *

from table1, table2;

ON \cup WHERE.

① Select _____

from A \rightarrow N

join B \rightarrow M

ON A.id = B.id

\Rightarrow A: [] , B: []

ans: []

for row1 in A :

for row2 in B:

if Cond is TRUE:

ans.add(row1, row2);

print(ans)

}

NxM.

⇒ Select —
from A

Cross Join B

Where A.id = B.id;

A: [] , B: []

temp: [] , final-ans: []

for row1 in A :

 for row2 in B:

 temp.add(row1, row2);

Cross
Join table

for row in ans :

 if condn is TRUE || WHERE.

 final-ans.add(temp[row])

⇒ 1st Approach is better.

————— * —————

Aggregate Queries.

- Aggregate functions \Rightarrow COUNT, MIN, MAX, SUM..
- Group By
- Having

Aggregate Queries.

Till now whatever queries we have written, they were returning some rows.

→ Find Students -----

→ Find batches -----

Q. Find the avg Rep of all the batches at scalar.

Output:

batch_id	avg rep
1	90
2	70
3	45

batches

batch_id	name	rep

Students

id	Name	batch_id	fee
1	A	1	90
2	B	1	80
3	C	1	70
4	D	Null	—
5	E	2	90

go through the Students table.



Combine rows of Students for each batch.



find the Avg batch.

Aggregate functions

↳ COUNT / SUM / Avg / MAX / MIN ...

→ COUNT() → 1 value.

Q. COUNT.

Students

id	name	per	batchid
1	A	80	1
2	B	70	1
3	C	90	Null
4	D	80	2

→ Find the count of students that have a batch.

COUNT: takes a lot of values in the input, combines them into a single value which is equal to count of non-null values.

COUNT(1, 2, 3, 4, 5) → 5

COUNT(1, 2, NULL, 4) → 3

Select COUNT(batchid)

for loop

[from Students.]



COUNT(1, 1, ~~NULL~~, 2) ⇒ 3
?

Print = Select COUNT(batch_id)
 from Students.
 where batch_id is NOT NULL;
 ↓
 COUNT(1, 1, 2) ⇒ 3

from
 ↓
 where
 ↓
Select
 = =

Note:- All aggregate functions ignore the NULL values.

Students

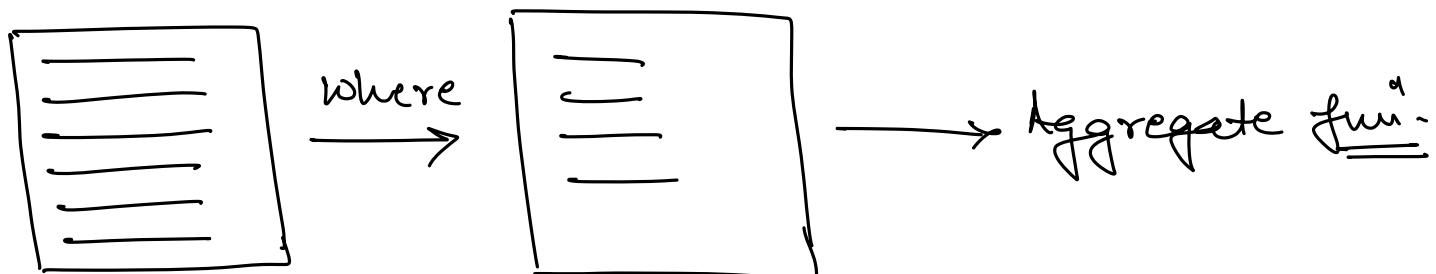
id	name	per	batch_id	age
1	A	80	1	22
2	B	70	1	21
3	C	90	Null	24
4	D	80	2	23

⇒ Count of all the students that have a batch and age <= 23.

Select COUNT(batch_id)
 from Students.
 Where age <= 23.

id	name	fee	batch_id	age
1	A	80	1	22
2	B	40	1	21

$\text{COUNT}(1, 1) \Rightarrow \underline{\underline{2}}$



Select $\underline{\underline{\quad}}$

from Students s
join batches b
ON s.batch_id = b.id;

→ Nested for loop.

from

Select.

Q.

Students

id	name	batch_id
1	A1	1
2	B1	1
3	C1	2

batches

id	name
1	A
2	B

get the count of Students with batch name R.

Select COUNT(s.id)

from Students S

join batches b

ON S.batch_id = b.id ;

Where b.name = 'A';

Intermediate table ans

Students			batches	
id	name	b_id	id	name
1	A1	1	1	A
2	B1	1	1	A
3	C1	2	2	B

JOIN

where.

ans2:

Students			batches	
id	name	b_id	id	name
1	A1	1	1	A
2	B1	1	1	A

⇒ Select count(id)
 from A
 join B
 ON Cond¹ ①
 Where Cond² ②

A: []

B: []

ans: []

```
for row1 in A:  

  for row2 in B:  

    if cond1 is TRUE:  

      ans.append(row1, row2)
```

ans2: []

JOIN.

for row in ans:
 if cond2 is TRUE:
 ans2.add(row); } Where.

Count_b = 0

for row in ans2:
 if row[id] != NULL:
 Count_b ++;

Print(Count_b)

$\Rightarrow \text{AVG}(1, 2, 3, \text{NULL}) \Rightarrow$

Aggregate
funⁿ.

$$\frac{1+2+3}{3} = \frac{6}{3} = 2$$

{ 2
3
Null } $\frac{6}{4}$

Students.

id	PSP
1	1
2	2
3	3
4	Null

$$\frac{\text{SUM(PSP)}}{\text{Count(id)}} \times \equiv \text{AVG(PSP)}$$

$$\rightarrow \frac{1+2+3}{3} = 2$$

* \Rightarrow Complete row.

Count(id)

Count(Rsp)

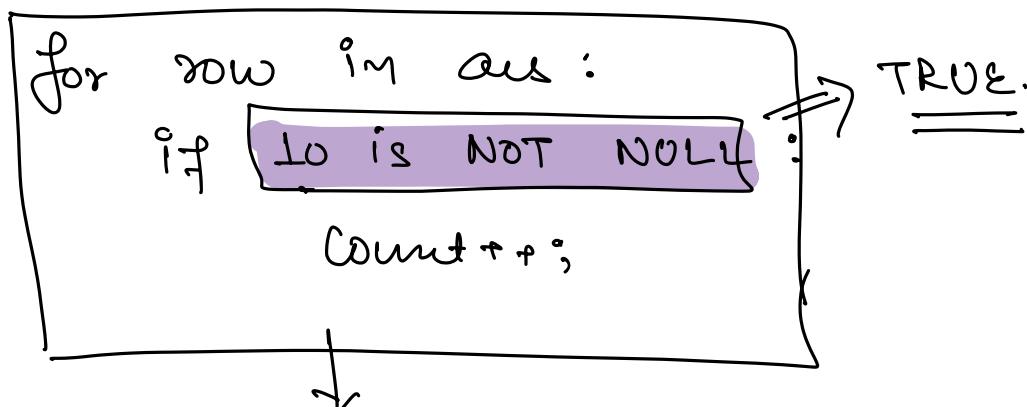
Count(b-id)

\Rightarrow COUNT(*)



COUNT (row₁, row₂, row₃, ...)

\Rightarrow Select Count(10) \Rightarrow DUMMY. (Always count all the rows).
from Students.



Count(*)

Efficient

Count(*) == Count(1)
Count(2)
Count('SCALER')
Count('UNITS')

<u>id</u>	<u>Name</u>	<u>batch</u>
<u>1</u>	<u>A</u>	<u>Null</u>
<u>Null</u>	<u>Null</u>	<u>Null</u>
<u>Null</u>	<u>Null</u>	<u>Null</u>

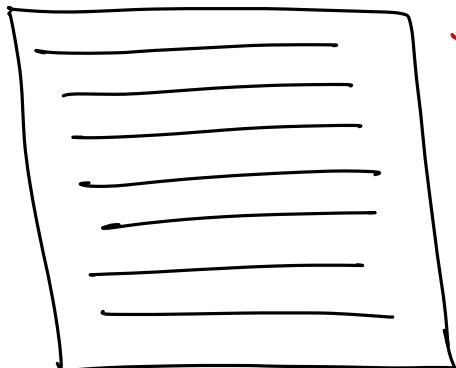
$\Rightarrow \text{Count(*)} \equiv \text{COUNT(1)}$

if (true) {

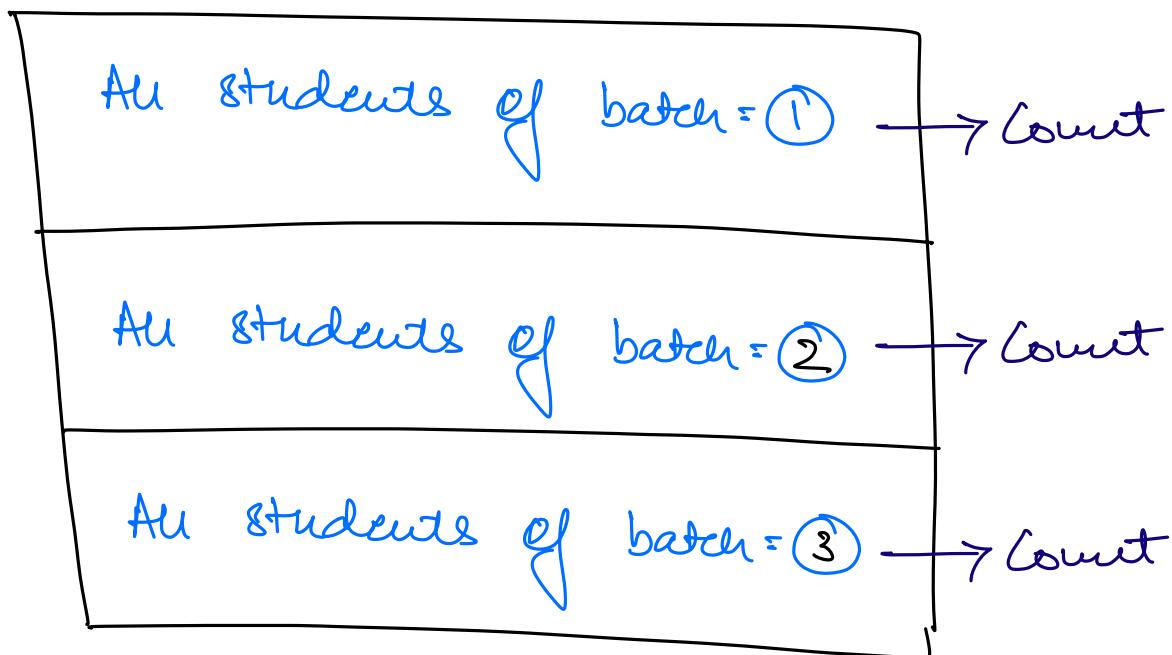
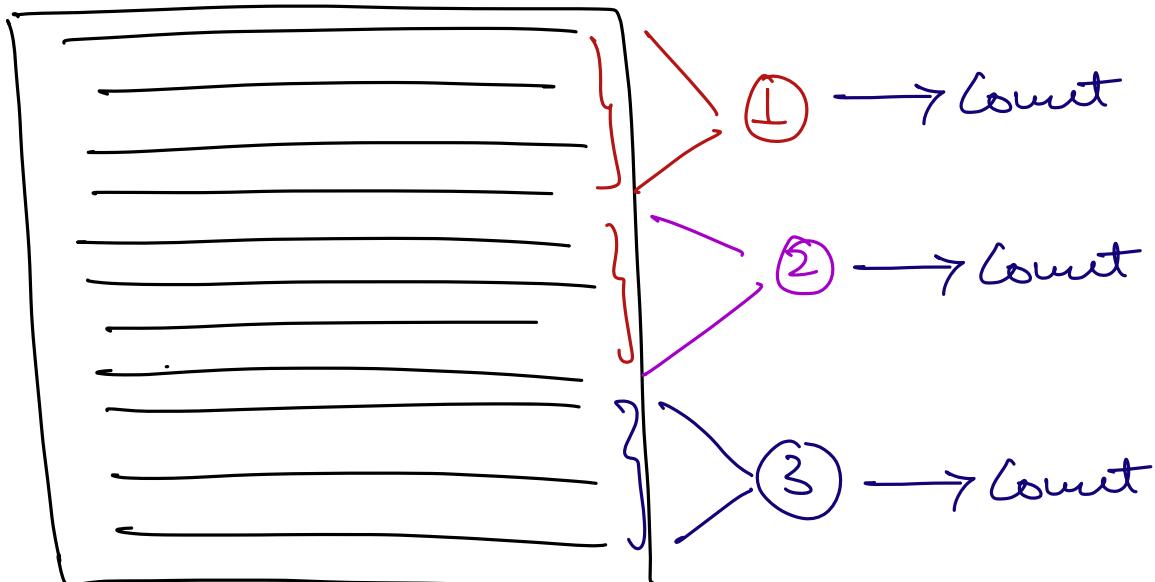
 3

Get the count of students for every batch.

b-id	count
1	50
2	70
3	40
4	60



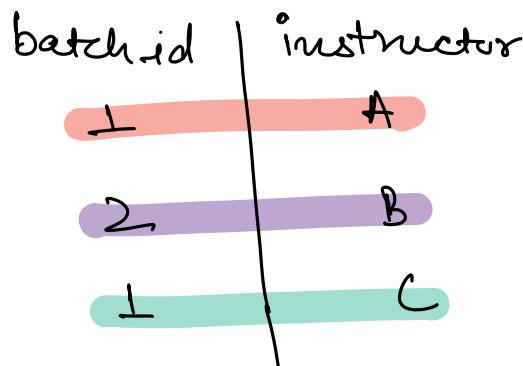
Count(*)



$\Rightarrow \underline{\text{GROUP}} \underline{\text{BY.}}$: Allows us to split the table into multiple groups and apply the aggregate funⁿ on groups

$\Rightarrow \text{group by batch_id}$
 ↳ brings all the Students of a batch together.

\Rightarrow group by batch_id, instructor ;



Count Students in each group.

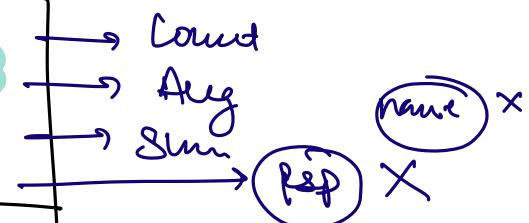
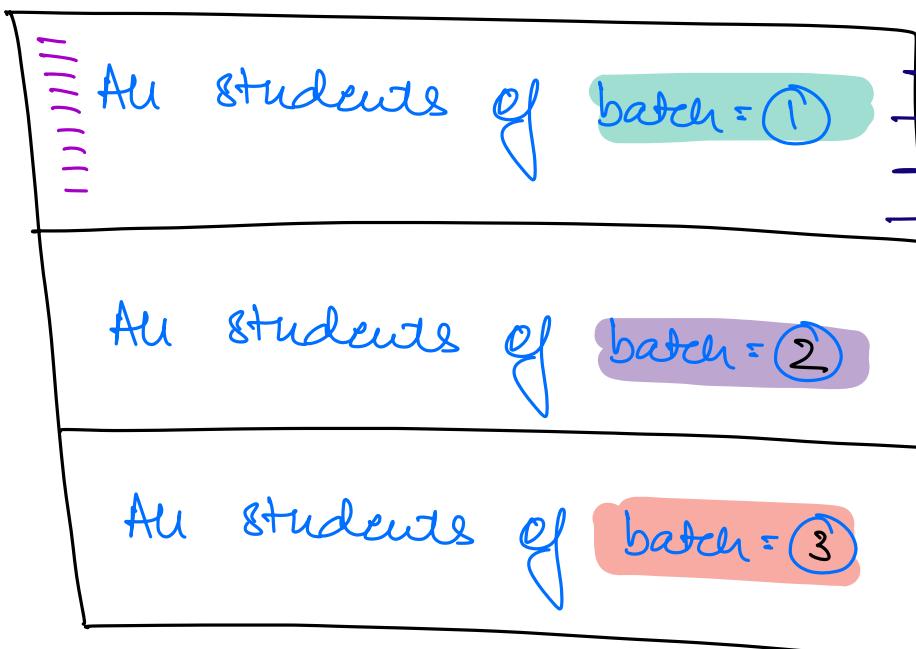
Select COUNT(1)

from student

group by batch_id

\Rightarrow

50
70
40
60



Select COUNT(1), batch_id

from student

group by batch_id

\Rightarrow

50	1
70	2
40	3
60	4

Select COUNT(1), batch_id, ~~P~~, ~~name~~, avg(Psp) ✓
 from student
 group by batch_id

⇒ If we are using group by clause then
 we can only use aggregate queries \ominus
 Columns in group by in Select statement.

Q. Print the batch names along with no. of
 students having more than 100 students.

Students

id	name	b_id

batches

id	name

Select batch_name, count(*)

from Students s

join batches b

ON s.b_id = b.id

group by b.name;

HAVING Count(*) > 100;

FROM
 ↓
 group by
 ↓
 HAVING
 ↓
 select.

temp-

student (s)			batch (b)	
id	name	b_id	id	name

batch names

(A)

(B)



temp

group (batch_name = A)	Count = 100
group (batch_name = B)	Count = 120
group (batch_name = C)	Count = 70

⇒ Where Clause :- Used to filter rows

⇒ filter groups : HAVING Clause.

⇒ After creating groups can we have where clause ?

No.

⇒ We can't add where clause after group by.

Q. Print the batch names along with no. of students which have $ps > 80$ having more than 100 students

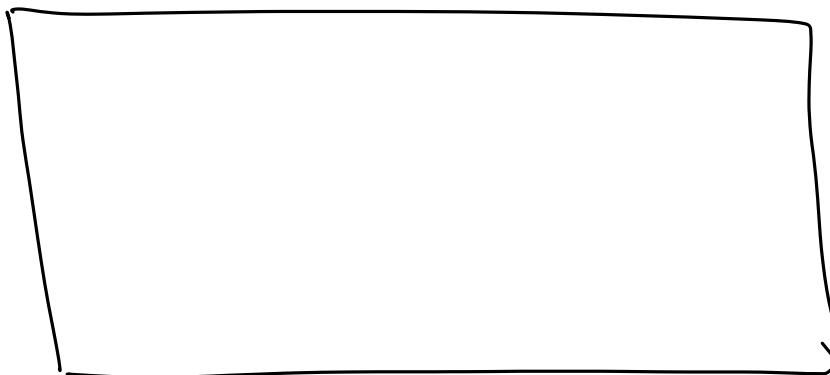
Select batch-name, count(*)
from Students S
join batches b
ON S.b_id = b.id
Where S.ps > 80.
group by b.name;
HAVING Count(*) > 100;

ans:

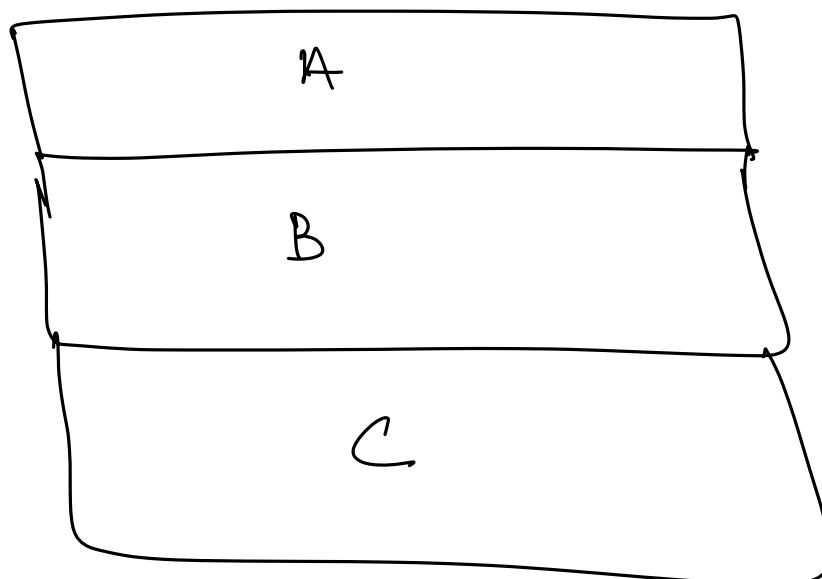
Students (S)			batches (b)	
id	name	b_id	id	name

↓ 80 (where)

Cens2



↓ Group by ~~Count~~.



↓ HAVING Count > 100

From

↓

Where

↓

Group by

↓

HAVING

↓

Select

Select -

— * —

Subqueries.

- Subqueries .
- IN .
- ALL
- ANY
- Correlated Subqueries .
- EXISTS .

Q:

Students

id	name	psp	batch_id

find all the students with psp > maximum psp of batch_id = 2 .

Students : [. . . .]

① find the MAX psp of batch_id = 2 ;

↳ \hat{x}

Select max(psp)
from student
where batch_id = 2

⇒ \hat{x}

II find all the students with $\text{PSP} \geq x$.

Select *
from Students
where PSP > x

Q: find all the students whose PSP >
PSP of student with id = 18.

Students (S1)

id	name	PSP

Students (S2)

id	name	PSP
18	XYZ	70

Select S1.id, S1.name
from Students S1
join Students S2
ON S2.id = 18 AND S1.PSP > S2.PSP

Self
Join.

- ① find rep of $\underline{\text{id}} = 18$. $\Rightarrow \text{X}$
- ② find students with $\underline{\text{rep}} > \text{X}$.

SUBQUERY: When we break a problem into smaller subproblems and use their ans to find the final ans.

\Rightarrow Most of the times problems that we solve via subqueries can also be solved by some other type of queries but generally subqueries are easier to understand.

① find rep of student with $\underline{\text{id}} = 18$. $\Rightarrow \text{X}$

② find students with $\underline{\text{rep}} > \text{X}$.

① Select rep
from students
where $\underline{\text{id}} = 18$; $\Rightarrow \text{X}$

② Select *
from students
where $\underline{\text{rep}} > \text{X}$.



Select *

from Students

Where Rep > (Select Rep

from Students
where id = 18)

→ N

⇒ Subqueries are easier to understand.

⇒ Subqueries :- Bad Performance.

TC : O(N²)

Above Subquery is returning a single no. (Rep of id=18).

Q. Find all the students with PSP > MAX PSP
of batch_id = 2

```
Select *
from Students
where PSP > (Select max(PSP)
                from Students
                where batch_id = 2)
```

Subquery is returning a
single value.

Subqueries can return multiple values as well

No. of rows	No. of columns	Output
1	1	Single value
1	M	Single row.
M	1	Single column
M	M	table.

>, <, >=, <=, = : Single value.

12.

Q.

Users

id	name	is_student	is_ta.
1	AB ✓	T	F
2	XY		F
3	Pranav	✗ T	F ✓
10	AB	F	(T)

⇒ Get the name of Students that are also

Name of TA's =

(St)
users

(ta)
users

id	name	is_student	is_ta.
1	Vishal	T	F/T
2	Deepak	F	

id	name	is_student	is_ta.
10	Vishal		(F)

Select _____

from Students St
join tas ta

ON St.name = ta.name ;

⇒

```

    select distinct(st.name)
    from users st
    join users ta
    ON st.is_student = true AND
        ta.is_ta = true AND
        st.name = ta.name.
  
```

- ① find name of all the TA's. ⇒ []
- ② find all the students with name ↗

① Select name
 from users ⇒ []
 where is_ta = true;

② Select distinct(name)
 from users
 where is_student = true AND
 name IN (Select name
 from users
 where is_ta = true);
 ['Deepak', 'Akashish', 'Pranav']

Q. Get the name of all the students whose PSP is not less than smallest rep of any batch.

batch_id	min_psp
1	40
2	30
3	50

50	✓
20	✗
32	✗
45	✗
52	✓
40	✓

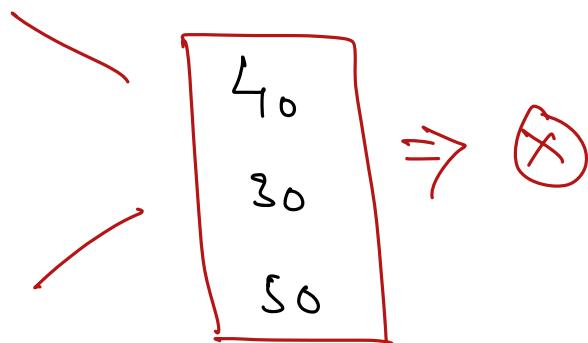
Not less than any value.

(OR)

greater than v, all values.
equal to

$n >= \text{Max of the min psp's of all the } \underline{\text{batches}}$.

① Select MIN(PSP)
from Students
group by batch_id



II Select MAX (PSP) }
 from Students X } \Rightarrow Y \Rightarrow SO

III Select *
 from Students
 where PSP >= Y

3 Select id, name
 from Students

where PSP > (Select MAX(min-PSP))

from (Select MIN(PSP) as min-PSP
 from Students
 group by batch_id) minPSPBatches

(2)) (1)

SUBQUERY inside FROM.

- ① We can give subquery inside from clause as well.
- ② Output of the subquery will be considered as a table.
- ③ We need to give a name to the subquery table if we are using with in from.

#

$n \geq \text{ALL}(10, 20, 30, 40)$

n should be greater than $10, 20, 30,$ & 40

v
or
equal to

ALL

$\Rightarrow \underline{\text{AND}}$

Select *

from students

where PSP $\geq \text{ALL}(\text{Select MIN(PSP)}$

from students

group by batch_id)

ANY. \Rightarrow OR.

Q. Get all the students with PSP greater than
avg PSP of their batch.

student, bid = 2 , PSP > Avg PSP of bid = 2

bid = 3 , PSP > Avg PSP of bid = 3 .

I Get Students with PSP > \times

II $X = \text{Avg PSP of student's batch}$

Select *

from Student st

where PSP > (Select avg(PSP)
from Students
where batch_id = st.batch_id)

\Rightarrow Correlated Subqueries.

$\Rightarrow St$
for student in Students:

for row in students

if (row.bid = st.bid) {

 aug-fep = _____

 3

 3

 if (st.fep > aug-fep) {

 // Add in the ans

 3

3
=

exists.

Q: Get all the students who are also TA.

Students

id	name	fep
1	Himanshu	_____

ta

id	name	st-id
—	—	10
—	—	Null
100	thu	1



Select st.id

from ta

where st.id IS NOT NULL;

✓

\Rightarrow select id
from Students

where id IN
IM
(Linear search)

TC: IM.

Select st_id
from ta
Where st_id IS NOT NULL;
[1, 10, 12, 15].

$O(N)$

select *

from Students \textcircled{S}

where EXISTS

Select st_id
from ta \textcircled{t}
Where t.st_id = g.id

This query will be
faster

#

Students

id	St-name
1	Pranjal ✓
2	

PK

mentor_sessions

Sid	s_id	mentorid
100	1	1
101	2	1
102	1	1
103	2	1
104	1	1
105	1	1
106	1	1
107	1	1
108	1	1
109	1	1
110	1	1
111	1	1
112	1	1
113	1	1
114	1	1
115	1	1
116	1	1
117	1	1
118	1	1
119	1	1
120	1	1
121	1	1
122	1	1
123	1	1
124	1	1
125	1	1
126	1	1
127	1	1
128	1	1
129	1	1
130	1	1
131	1	1
132	1	1
133	1	1
134	1	1
135	1	1
136	1	1
137	1	1
138	1	1
139	1	1
140	1	1
141	1	1
142	1	1
143	1	1
144	1	1
145	1	1
146	1	1
147	1	1
148	1	1
149	1	1
150	1	1
151	1	1
152	1	1
153	1	1
154	1	1
155	1	1
156	1	1
157	1	1
158	1	1
159	1	1
160	1	1
161	1	1
162	1	1
163	1	1
164	1	1
165	1	1
166	1	1
167	1	1
168	1	1
169	1	1
170	1	1
171	1	1
172	1	1
173	1	1
174	1	1
175	1	1
176	1	1
177	1	1
178	1	1
179	1	1
180	1	1
181	1	1
182	1	1
183	1	1
184	1	1
185	1	1
186	1	1
187	1	1
188	1	1
189	1	1
190	1	1
191	1	1
192	1	1
193	1	1
194	1	1
195	1	1
196	1	1
197	1	1
198	1	1
199	1	1
200	1	1
201	1	1
202	1	1
203	1	1
204	1	1
205	1	1
206	1	1
207	1	1
208	1	1
209	1	1
210	1	1
211	1	1
212	1	1
213	1	1
214	1	1
215	1	1
216	1	1
217	1	1
218	1	1
219	1	1
220	1	1
221	1	1
222	1	1
223	1	1
224	1	1
225	1	1
226	1	1
227	1	1
228	1	1
229	1	1
230	1	1
231	1	1
232	1	1
233	1	1
234	1	1
235	1	1
236	1	1
237	1	1
238	1	1
239	1	1
240	1	1
241	1	1
242	1	1
243	1	1
244	1	1
245	1	1
246	1	1
247	1	1
248	1	1
249	1	1
250	1	1
251	1	1
252	1	1
253	1	1
254	1	1
255	1	1
256	1	1
257	1	1
258	1	1
259	1	1
260	1	1
261	1	1
262	1	1
263	1	1
264	1	1
265	1	1
266	1	1
267	1	1
268	1	1
269	1	1
270	1	1
271	1	1
272	1	1
273	1	1
274	1	1
275	1	1
276	1	1
277	1	1
278	1	1
279	1	1
280	1	1
281	1	1
282	1	1
283	1	1
284	1	1
285	1	1
286	1	1
287	1	1
288	1	1
289	1	1
290	1	1
291	1	1
292	1	1
293	1	1
294	1	1
295	1	1
296	1	1
297	1	1
298	1	1
299	1	1
300	1	1
301	1	1
302	1	1
303	1	1
304	1	1
305	1	1
306	1	1
307	1	1
308	1	1
309	1	1
310	1	1
311	1	1
312	1	1
313	1	1
314	1	1
315	1	1
316	1	1
317	1	1
318	1	1
319	1	1
320	1	1
321	1	1
322	1	1
323	1	1
324	1	1
325	1	1
326	1	1
327	1	1
328	1	1
329	1	1
330	1	1
331	1	1
332	1	1
333	1	1
334	1	1
335	1	1
336	1	1
337	1	1
338	1	1
339	1	1
340	1	1
341	1	1
342	1	1
343	1	1
344	1	1
345	1	1
346	1	1
347	1	1
348	1	1
349	1	1
350	1	1
351	1	1
352	1	1
353	1	1
354	1	1
355	1	1
356	1	1
357	1	1
358	1	1
359	1	1
360	1	1
361	1	1
362	1	1
363	1	1
364	1	1
365	1	1
366	1	1
367	1	1
368	1	1
369	1	1
370	1	1
371	1	1
372	1	1
373	1	1
374	1	1
375	1	1
376	1	1
377	1	1
378	1	1
379	1	1
380	1	1
381	1	1
382	1	1
383	1	1
384	1	1
385	1	1
386	1	1
387	1	1
388	1	1
389	1	1
390	1	1
391	1	1
392	1	1
393	1	1
394	1	1
395	1	1
396	1	1
397	1	1
398	1	1
399	1	1
400	1	1
401	1	1
402	1	1
403	1	1
404	1	1
405	1	1
406	1	1
407	1	1
408	1	1
409	1	1
410	1	1
411	1	1
412	1	1
413	1	1
414	1	1
415	1	1
416	1	1
417	1	1
418	1	1
419	1	1
420	1	1
421	1	1
422	1	1
423	1	1
424	1	1
425	1	1
426	1	1
427	1	1
428	1	1
429	1	1
430	1	1
431	1	1
432	1	1
433	1	1
434	1	1
435	1	1
436	1	1
437	1	1
438	1	1
439	1	1
440	1	1
441	1	1
442	1	1
443	1	1
444	1	1
445	1	1
446	1	1
447	1	1
448	1	1
449	1	1
450	1	1
451	1	1
452	1	1
453	1	1
454	1	1
455	1	1
456	1	1
457	1	1
458	1	1
459	1	1
460	1	1
461	1	1
462	1	1
463	1	1
464	1	1
465	1	1
466	1	1
467	1	1
468	1	1
469	1	1
470	1	1
471	1	1
472	1	1
473	1	1
474	1	1
475	1	1
476	1	1
477	1	1
478	1	1
479	1	1
480	1	1
481	1	1
482	1	1
483	1	1
484	1	1
485	1	1
486	1	1
487	1	1
488	1	1
489	1	1
490	1	1
491	1	1
492	1	1
493	1	1
494	1	1
495	1	1
496	1	1
497	1	1
498	1	1
499	1	1
500	1	1

→ Get all the students names those have taken a mentor session.

I.M.

Select *
from Students
where EXISTS
(Select *
from mentor_sessions ms
where ms.s_id = s.id)

↓

{ get at least one ∞ ,
then it won't execute
further. }

IN $\textcircled{v.s.}$ EXISTS.

Agenda:

- Indexes.
- How indexes work?
- Indexes for range queries.
- Cons of Indexes
- How to create indexes?

⇒ Indexing.

[for $i \rightarrow 1$ to N
[for $j \rightarrow 1$ to M
ans.add (row)]] } $O(NM) \Rightarrow$

1B rows \Rightarrow 10¹⁸

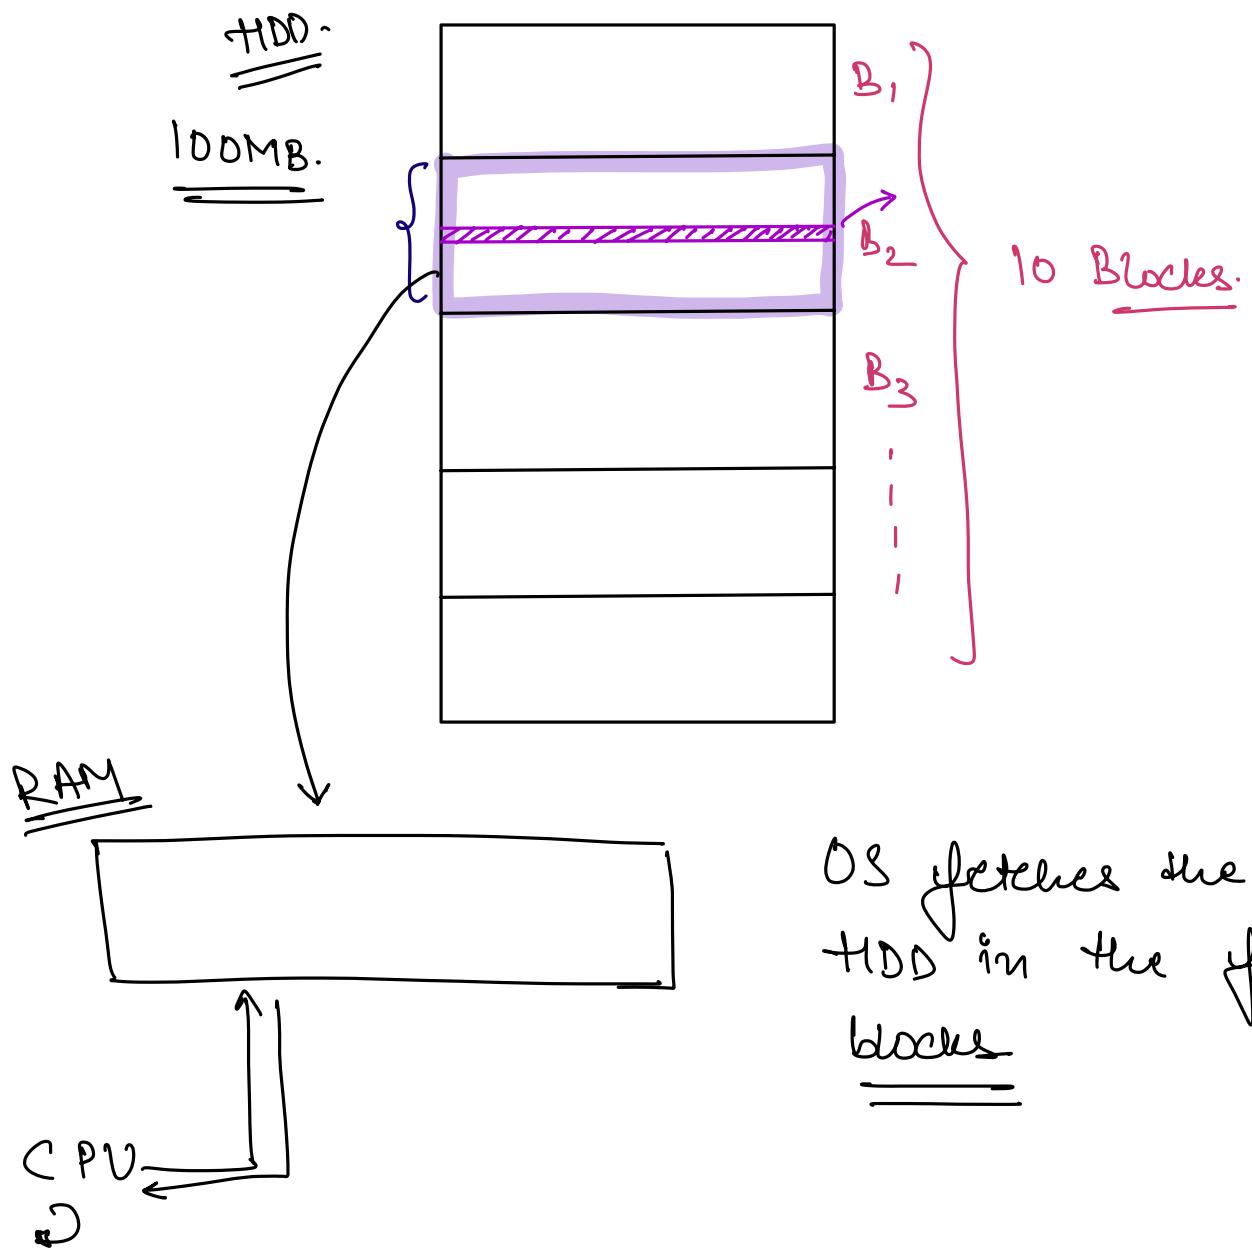
⇒ Performance will be very very slow.

① DB needs to do lot optimizations to make our queries faster.

② DB stores the data on disk.

⇒ CPU can't fetch data directly from disk.

⇒ CPU can fetch data from RAM.

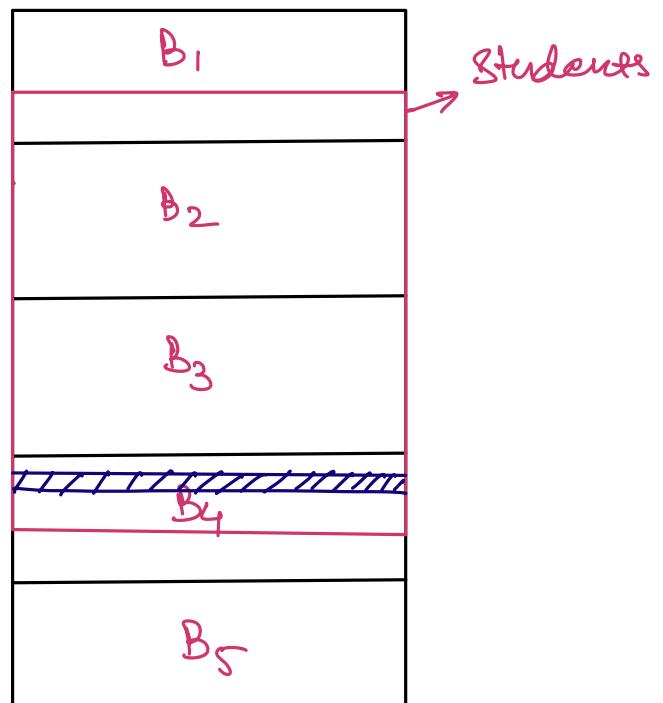


OS fetches the data from
HDD in the form of
blocks

⇒

Select *
from Students
where id = 100

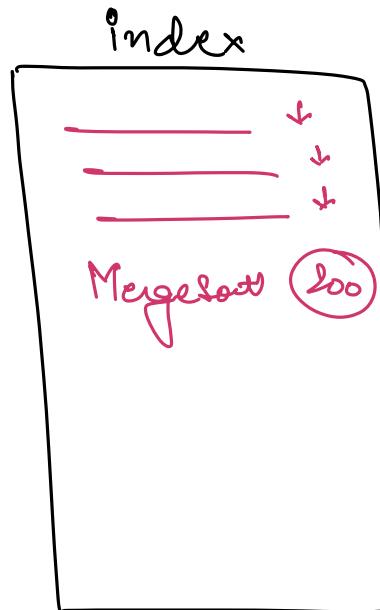
- ① Bring B₁ to the RAM:
- ② Check if id = 100 is present in B₁ B₂ B₃



Time ⇒ # of block accesses from disk.

If we want to optimise the overall performance of DB, then we need to optimise # of block accesses from HDD.

INDEX.



CLRS. ≈ 1000 Pages

Index table can reduce no. of false page accesses.

index in book \Rightarrow find pages faster.

index in DB \Rightarrow Optimises the # of block accesses.

Purpose of Indexing : Reduces the disk accesses

⇒ How Indexes work ?

Table : 100M rows.

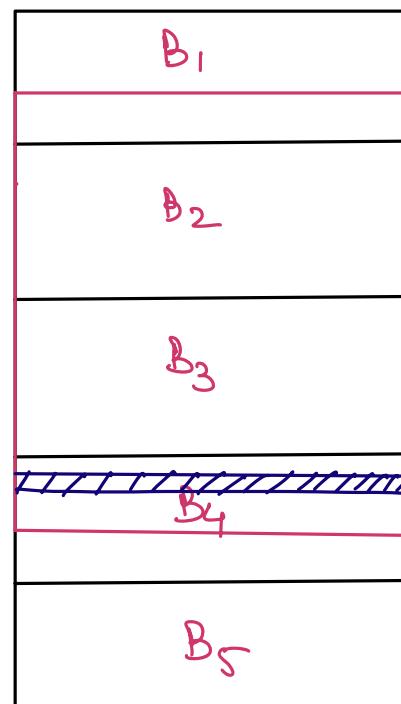
Select *
from Students
where Id = 100

Index (HashMap | Map)

table

id	block
1	B ₁
2	B ₁
3	B ₂
4	B ₂
5	B ₃
...	...
100	B ₄

RAM
&
HDD



\Rightarrow TC: O(1) Avg.

```
# Select *
from Students
where name = 'Abhishek';
```

Index \Rightarrow HashMap.

name	List<block>
'Abhishek'	<1, 2, 3, 5, 10, 18>
'Priya'	<5, 6>
'Soham'	<2, 4, 10>
'Deepak'	<6, 9>
==	==
==	==

without index \Rightarrow All the blocks.

with index \Rightarrow only 6 blocks.

Index table \equiv HashMap

\downarrow
<key, value>

Select *
 from Students
 where PSP between 40 AND 80;
double.

40	B ₁
45	B ₂
55	B ₃
41	:
:	:
60	
65	
59	
52	
47	

40.391 50.86
 70.68

Range Queries.

HashMap<Int, String>

~~O(NP)~~
 for i >= 40 & i <= 80;
 if (i in HM) {
 map.get(i);

HashMap / unordered-map.

$40 \leq PSP \leq 80 \Rightarrow$ Ordering

↳ TreeMap | ordered-map.

Balanced BST.

Height : $O(\log n)$

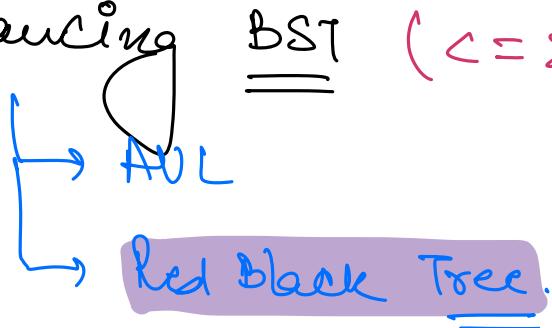
Always

$10^6 \Rightarrow O(1)$

$\Rightarrow \log 10^6 \approx 20$

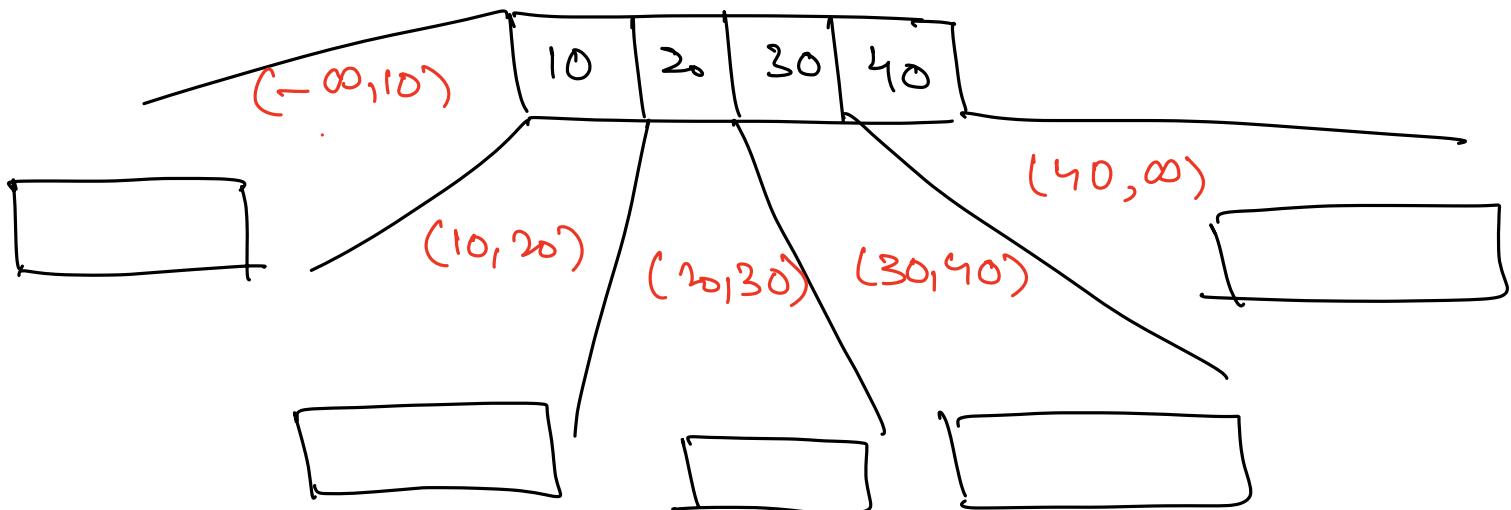
- I) Indexes work very similarly to TreeMaps.
- II) Index: DS that reduces the no. of disk accesses to optimise the queries.

\Rightarrow Self Balancing BST (≤ 2 children)



$\Rightarrow \log N$

\Rightarrow B | B+ Trees.
($\leq x$ children)

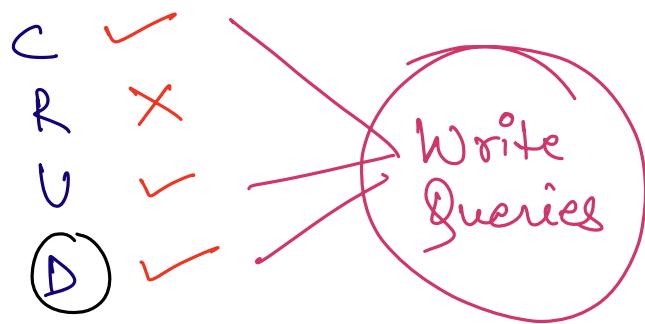


$H < \log N$.

CONS of indexing.

Index table

-	-



- ⇒ Writes becomes slower.
- ⇒ Storage also increases.

Note: Don't create indices blindly.

Create an index only when we start getting frequent queries on any table for a column.

Index of Multiple columns.

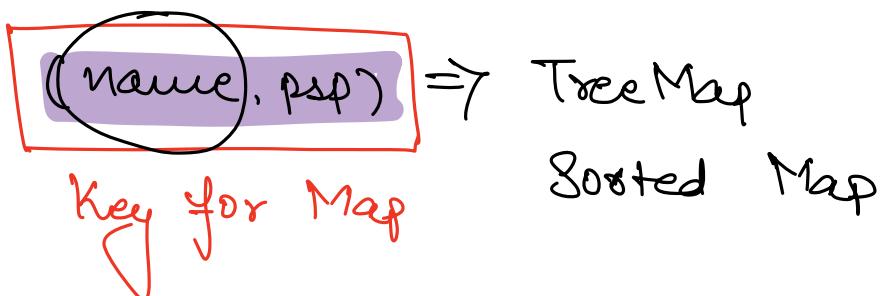
(id)

(name)

(name, psp)

I) (name, psp)

II) (psp, name)



(name, psp)	Block
('Abhishek', 80)	1
('Abhishek', 90)	1
('Abhishek', 92)	2
(Ajay, 20)	3
(Bijoy, 80)	1
(Deepak, 60)	1
=====	1
(____) 80	1

Select *
from Students
where name = 'Abhishek'
=====

Select *
from Students
where psp = 80;

X

Index on (name, psp) won't optimise queries
based on psp.

(PSP, name)

(PSP, name)	Block
10, —	
10, —	
15, —	
20, —	
40, —	

Select *
from Students
Where PSP = 80; ✓

Select *
from Students

Where name = 'Abhishek' AND
PSP = 80;

⇒ (name, PSP)

SOO
==

⇒ Index | Query | Help / Not Help.

name

PSP

X

(name, PSP)

name

✓

(name, PSP, attendence)

name

✓

(name, PSP, attendence)

(name, PSP)

✓

(name, sp, attended)

(rep)

X

If a query is on column ~~X~~ any index with prefix ~~X~~ would optimise the query

⇒ Index on Strings.

```
Select *  
from Students  
Where email = —
```

1M
Entries

email	block
deepak@ Scaler.com	1
tinausteen@ gmai1.com	2
80ham@ yahoo.com	3
—	
—	
—	

B|B+ Trees

_____ @ —

email	block
deepak	[2, 8, 5, 4]
timashu	[1, 4]
Soham	[2, 3]
gaffar.	:
:	:

⇒ < 1M entries.

⇒ Space vs Optimization.

⇒ name

Select *

from Students

Where name = _____;

⇒ Index on name Column

⇒ Ideally when we have to create index on ~~strings~~ column, it's better to create index only on first few characters.

Prefix

⇒ first 5 characters

⇒ 2B users.

index on how many characters	# of users	# of unique entries	No. of Block accesses.
1	2×10^9	<u>26</u>	<u>$2 \times 10^9 / 26$</u> (Avg)
2	2×10^9	26×26	$2 \times 10^9 / 26 \times 26$
3	2×10^9	$26 \times 26 \times 26$	$2 \times 10^9 / 26 \times 26 \times 26$
:	:	:	:

ab → 2B | 26x26

ac →

ad →

:

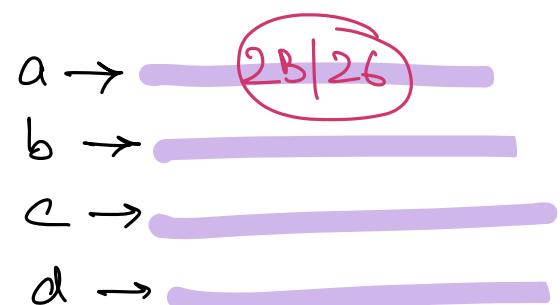
:

abc

abd

abe

|



Agenda :

- Transactions.
- Properties of Transactions :

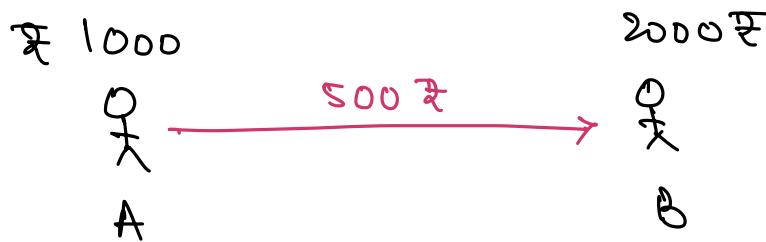
A
C
I
D

→ Commit & Rollback

→ Isolation levels.

⇒ What are Transactions ?

Bank Server.

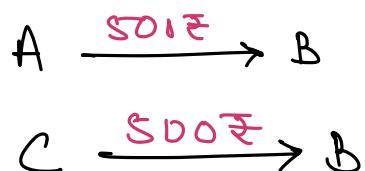


accounts.

id	name	balance	type
1	A	1000 <small>500</small>	Savings	
2	B	2000 <small>2500</small>	Saving	

Transfer Money :

- i) Get balance of A from DB \Rightarrow Select
- ii) Check balance $\geq 500 \Rightarrow$ Application.
- iii) Reduce balance of A by 500 \Rightarrow Update
- IV) Add 500 \Rightarrow B \Rightarrow Update



Transfer Money operation might be getting executed for multiple users at the same time.

$$\boxed{n += 10} \Rightarrow \text{Not a single op"}$$

$$n = n + 10.$$

```
temp  $\leftarrow$  n  
temp  $\leftarrow$  temp + 10  
n  $\leftarrow$  temp.
```

A:	1000	500
B:	5000	5500 6000
C:	4000	3500

}

$$A \xrightarrow{500} B$$

1000

- ① $X \leftarrow \text{read } A$
- ② $X \leftarrow X - 500$
- ③ $A \leftarrow X$
- ④ $X \leftarrow B$
- ⑤ $\underline{500}$
- ⑥ $X \leftarrow X + 500$
- ⑦ $B \leftarrow X$

$$C \xrightarrow{500\%} B$$

4000

- ⑧ $X \leftarrow \text{read } C$
- ⑨ $X \leftarrow X - 500$
- ⑩ $C \leftarrow X$
- ⑪ $Y \leftarrow \boxed{B}$
- ⑫ $Y \leftarrow Y + 500$
- ⑬ $B \leftarrow Y$

$$A \xrightarrow{500} B$$

1000

- ① $X \leftarrow \text{read } A$
- ② $X \leftarrow X - 500$
- ③ $A \leftarrow X$
- ④ $X \leftarrow B$
- ⑤ $X \leftarrow X + 500$
- ⑥ $B \leftarrow X$

$$\begin{aligned} A &: \cancel{1000} \ 500 \\ B &: \cancel{500} \ \cancel{500} \ \cancel{500} \\ C &: \cancel{4000} \ 3500 \end{aligned} \quad \left. \begin{array}{l} \text{Inconsistent} \\ \hline \end{array} \right.$$

$$C \xrightarrow{500\%} B$$

4000

- ⑦ $Y \leftarrow \text{read } C$
- ⑧ $Y \leftarrow Y - 500$
- ⑨ $C \leftarrow Y$
- ⑩ $Y \leftarrow B$
- ⑪ $Y \leftarrow Y + 500$
- ⑫ $B \leftarrow Y$

\Rightarrow If multiple operations happens on same data at the same time it can lead to inconsistent.

Transaction: Set of DB operations logically grouped together to perform a task.

transferMoney:-

- i) Get balance of A from DB \Rightarrow Select
- ii) Check balance $\geq 500 \Rightarrow$ Application.
- iii) Reduce balance of A by 500 \Rightarrow Update
- iv) Add 500 \Rightarrow Update

transaction.

\Rightarrow Expectations / Properties of a Transaction:

A : Atomicity

C : Consistency

I : Isolation

D : Durability

ATOMICITY.



Smallest Indivisible unit.

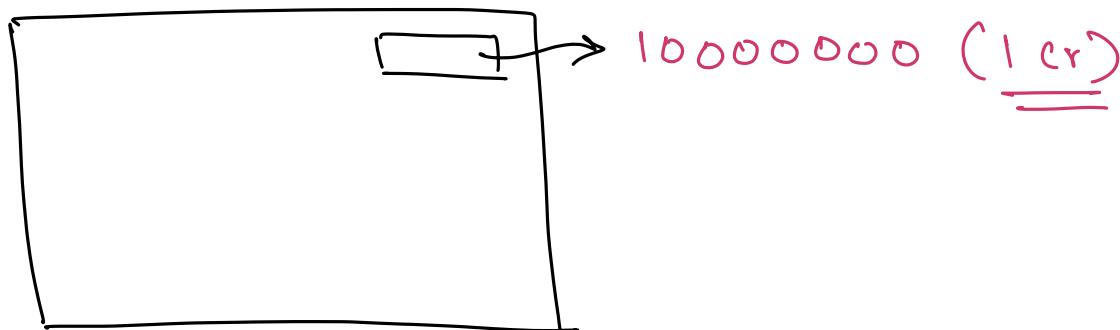
- ⇒ Transaction is the smallest unit.
- ⇒ To the end user transaction should appear single indivisible operation.
- ⇒ Either the complete transaction should happen or nothing should happen.
- ⇒ Any transaction should not stop at any intermediary step.

CONSISTENCY.

↳ Data correctness.

Database should be in consistent state before & after the transaction.

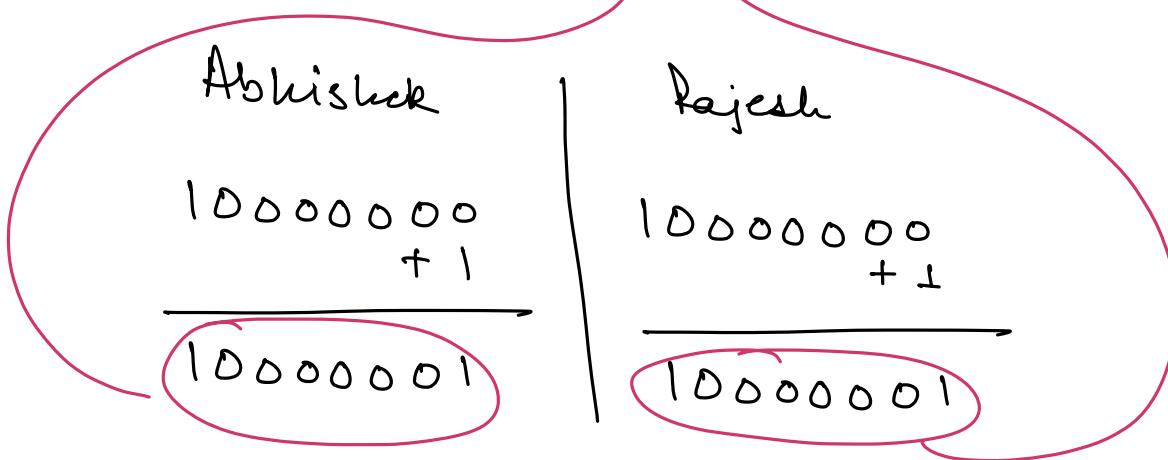
floatstar.



Abhishek
↓
audience

Rajesh
↓

	id	Count
	123	100000000 10000001



35000000 ⇒

35000100 ⇒

ISOLATION.

One transaction should not affect the execution of other transactions running at the same time on same data.

A \xrightarrow{x} B $\Rightarrow T_1$

C → B $\Rightarrow T_2$

} isolated.

$$\begin{array}{l} A \rightarrow B \Rightarrow T_1 \\ C \rightarrow D \Rightarrow T_2 \end{array} \quad \left. \right\}$$

DURABILITY.

↳ Persistence.

- ⇒ Once a transaction has completed we would want its changes to persisted in the DB permanently,
- ⇒ Write the changes to disk.

Commit & Rollback.

Students.

id	name	psp	- - -
1	Deepak	70	

update students
set psp = 80
where id = 1;

⇒ Select * from students
where id = 1; $\Rightarrow 80$

AutoCommit → 0
→ 1

Every SQL query starts a transaction by default and it auto commits the changes.

Starts query
executes
Commit

↳ We'll have to Commit the changes done by a transaction on our own.

Keyword.

Transaction Isolation Levels.

⇒ Most relaxed

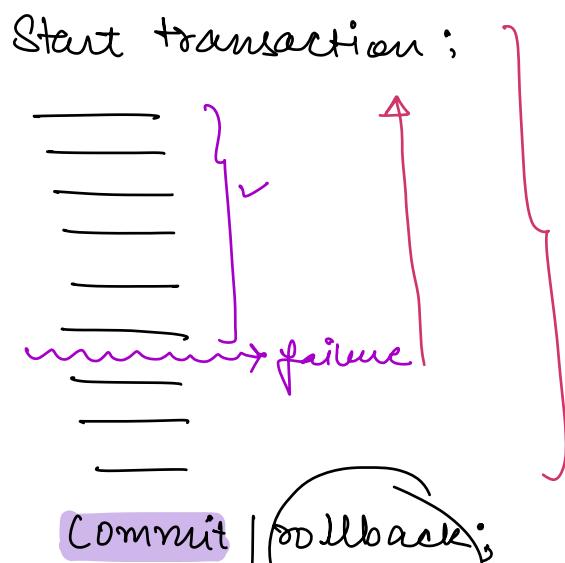
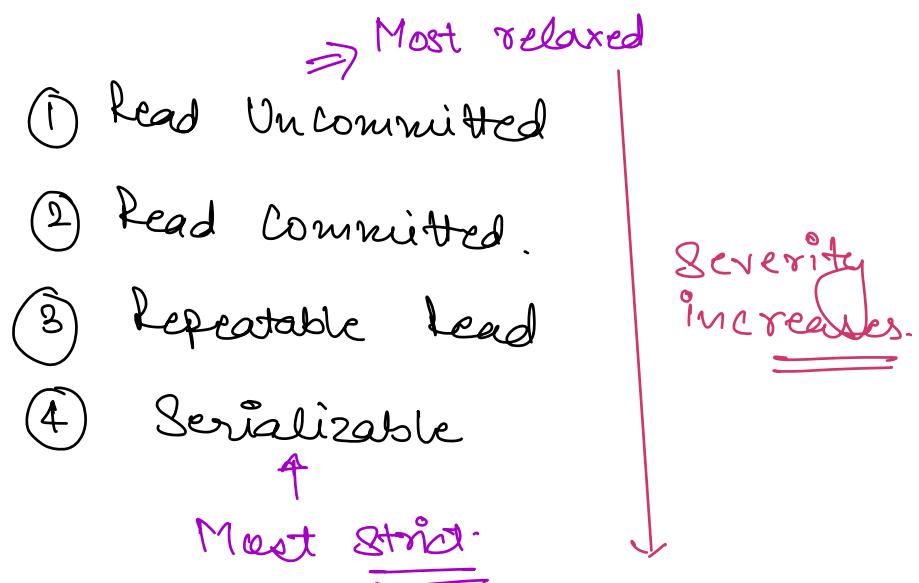
- ① Read Uncommitted
- ② Read Committed.
- ③ Repeatable Read
- ④ Serializable

↑
Most strict.

Severity increases.



Transaction Isolation Levels.

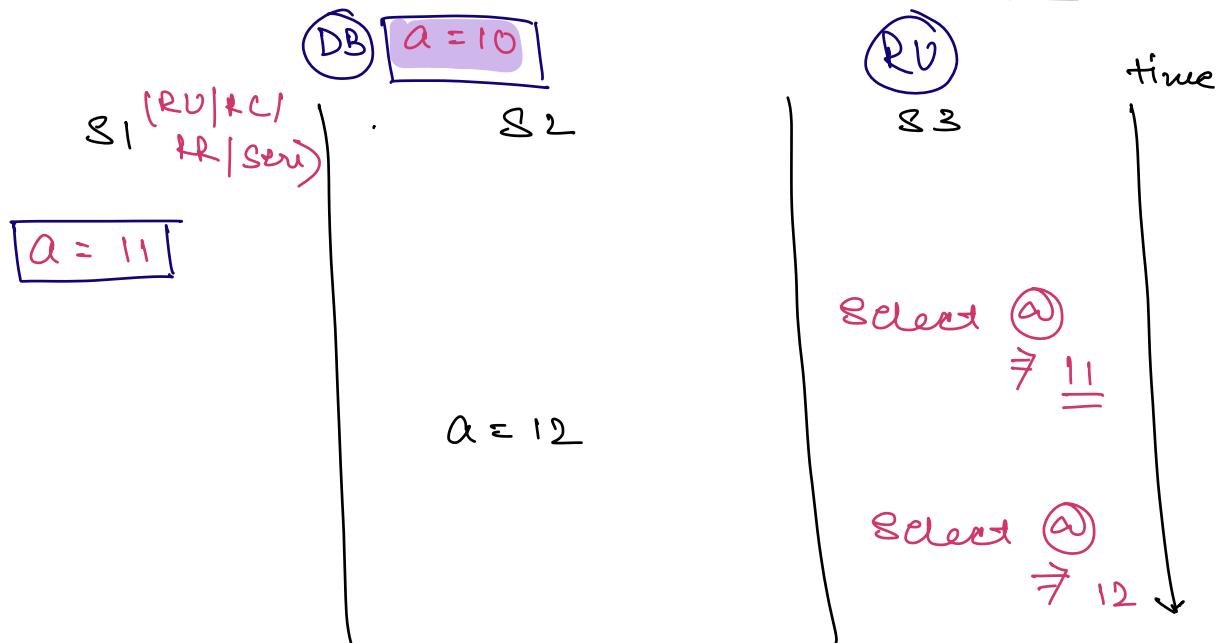


Commit | Rollback;

- ① By default autoCommit = true
- ② If we start a transaction then we'll have to commit on our own.

Read Uncommitted :-

⇒ Allows us to read the latest uncommitted.



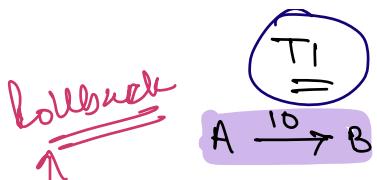
Note: The value that we are going to read will depend on our isolation level & not on other's isolation level.

⇒ MySQL has Repeatable Read isolation level.

A ◎ I D.
↓

$$A = \cancel{100} 140$$

$$B = \cancel{200} 150$$



- 1) $X \leftarrow \text{Read } A (100)$
- 2) $X \leftarrow X - 10 (90)$
- 3) $A \leftarrow X (90)$

~~8) $X \leftarrow \text{Read } A$~~ $\xrightarrow{\text{Edit}} X \leftarrow X + 10$

$B \leftarrow X$

COMMIT;



- 4) $X \leftarrow \text{Read } B (200)$
- 5) $X \leftarrow X - 50 (150)$
- 6) $B \leftarrow X (150)$

\downarrow
7) $X \leftarrow \text{Read } A (90)$

8) $X \leftarrow X + 50 (140)$

9) $A \leftarrow X (140)$

COMMIT;



Inconsistent State
of data.

Cause

$\Rightarrow \underline{\text{DIRTY READ}}$.

It can lead to inconsistent | wrong DB state
if we read recommitted data.

Prob.: fast

Read Committed.

- ⇒ Transaction will always read the committed data from disk.
- ⇒ No dirty read problem.
- ⇒ Stricter than Read Uncommitted.

Students

id	name	Psp	VOUCHER_SENT
1	A	82 ⁷⁹	false
2	B	75	false
3	C	85	false
4	D	40	false
5	E	60	false

Q. Send Voucher to all the students with $Psp > 80$.

list <Students> ⇒ Select *
from Students
where Psp > 80;

Vouchers Send. ⇒ 20 mins

{
 Update Students
 set Voucher-Sent = true
 where $\text{Rep} \geq 80$;

Read

Committed. \Rightarrow Student (A) won't come
in the output of T2
because $\text{Rep} < 80$.

\Rightarrow At both the places we should read the
same data.

\Rightarrow We should read the Repeated Data.

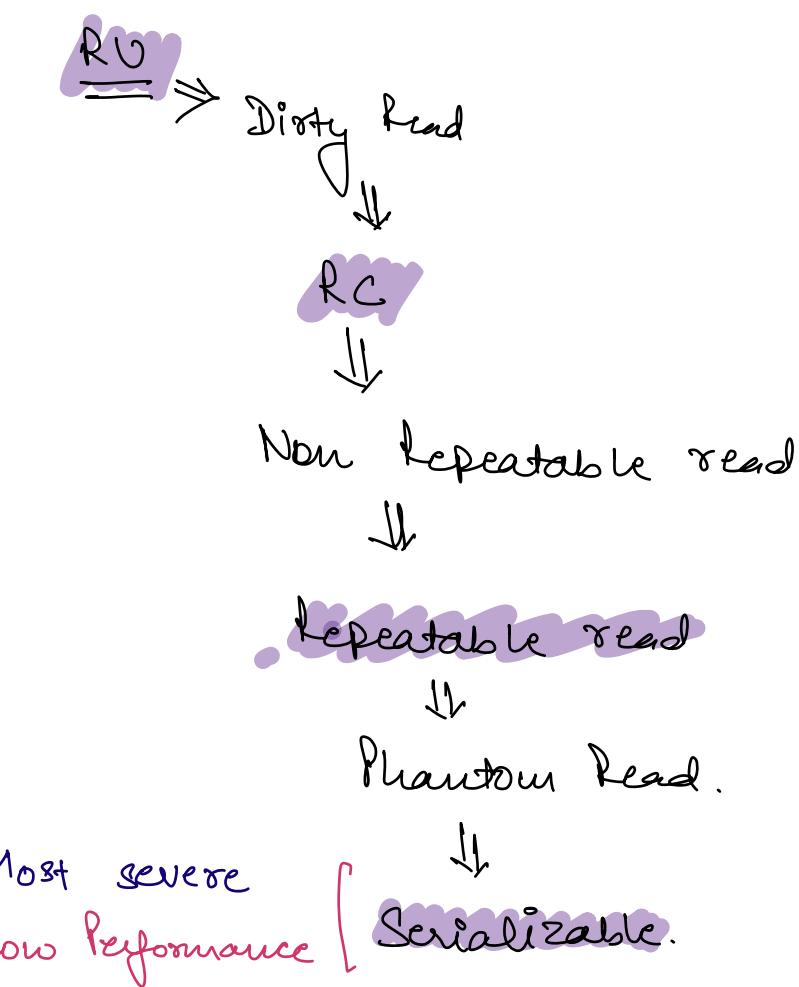
\Rightarrow Problem : Non Repeatable Read.

Repeatable Read.

⇒ In a transaction once we have read that we should read the same data again

Problem.

PHANTOM READ. : A row came magically from nowhere at the time of Update.



Serializable.

Repeatable Read + Checks for locks even for read operation.

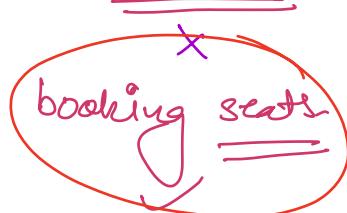
(If lock is present while reading then others won't be even able to read the data).



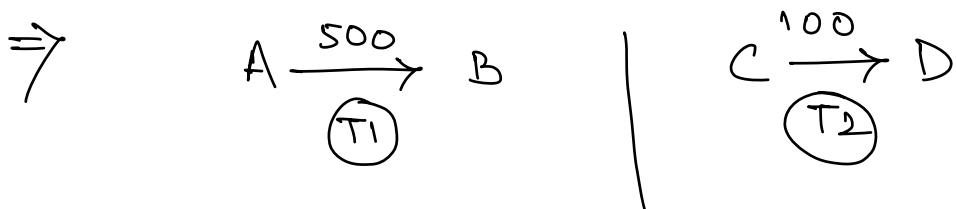
V
D

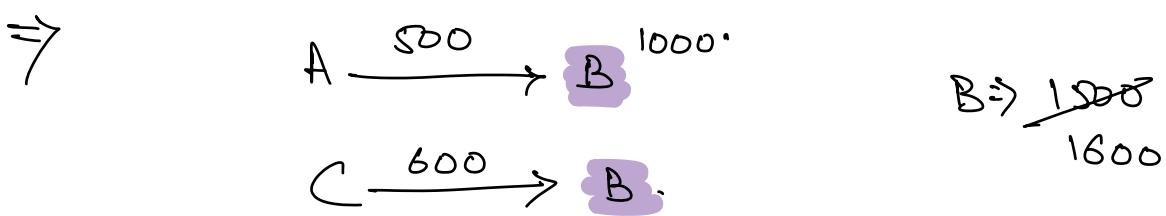
X	X	X	V	V	V	X

Show Seats

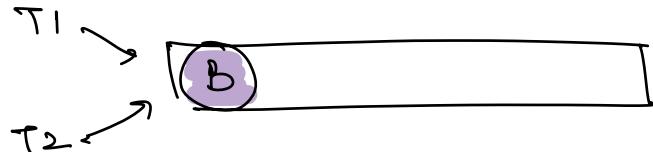


Booking -





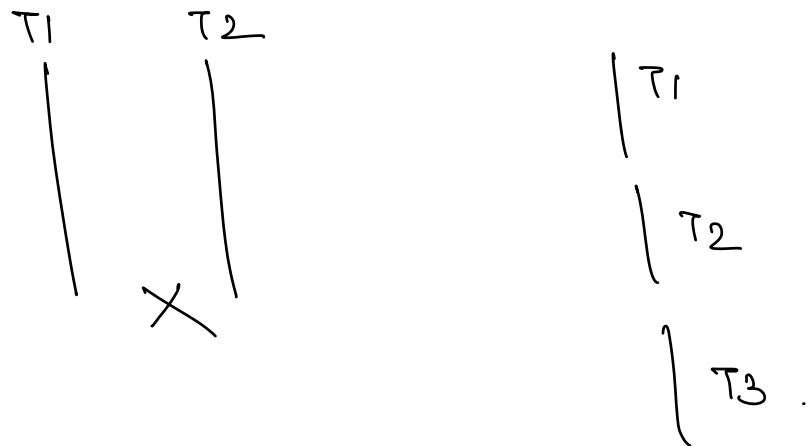
Serializable



$$B = 1000$$

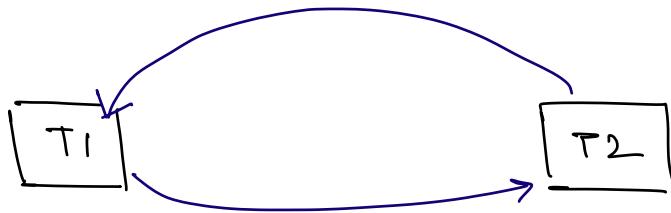
(T1) $A \xrightarrow{500} B \Rightarrow B = 1500$

(T2) $C \xrightarrow{600} B \Rightarrow B = \underline{\underline{2100}}$.



\Rightarrow One transaction can only happen on any set of data any given point of time.

#



- Lock on id = 1
- Trying to read id = 2
- Trying to read id = 1
- lock on id = 2

\Rightarrow Deadlock.

\Rightarrow MySQL will handle deadlocks by cancelling one of the transaction.

T1

Select *
from —
where id = 1
for update;

T2

Select *
from —
where id = 1

Select *
from —
where id = 2

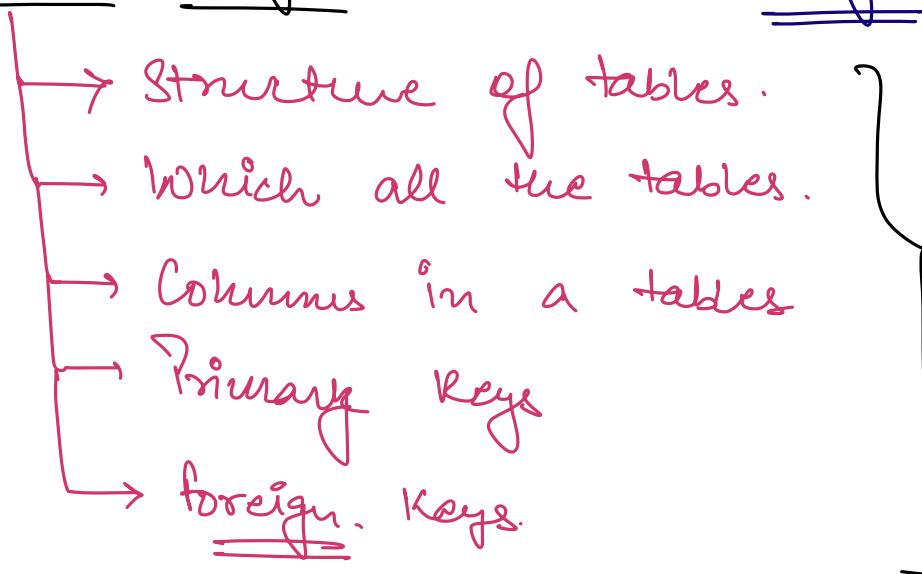
Select *
from —
where id = 2
for update;

Deadlock

Agenda:

1. What is Schema Design.
2. How to approach Schema Design.
3. Cardinalities
4. How to find cardinalities.

Schema Design ≡ Database Design.



⇒ Benefits of Schema Design

- Performance
- Reduce redundancy.
- Reduce Inconsistency

How to Approach Schema Design.

SCALER.

1. Scaler will have multiple **batches**, for each batch we'll store its name, start month and current instructor.
2. Each batch will have multiple students.
3. Each batch will have multiple **classes**.
4. For each class we need to store their name, date & time, instructor.
5. For each student we'll store their name, email, phone, grad year, university name, etc.
6. Every student will have a **buddy**, who is another student.
7. A student can be moved from one batch to another batch.
8. For each batch a student goes to we'll store the start date of the batch.
9. Every student can choose their **mentor**, for mentor we'll store their name, email & current company.
10. We'll store the information regarding **mentor sessions** like session-id, time, duration, student, mentor, student rating, mentor rating

ii. for every batch we'll store if it is Academy
& DSML

STEPS.

① Create the tables.

⇒ how to identify tables.

→ find out all the Nouns in the requirements
& check if we need to store data.

Conventions about names

1) Table names should be plural.

Students

Mentors

Batches

Mentor-Sessions

MentorSessions

2) Column name ⇒ singular.

3) Primary Keys.

Batches

id / batch_id

Students .

batches

	batch-id		
--	----------	--	--

Using .

Tables .

batches

batch-id	name	start month	
----------	------	-------------	--

instructors

instructor-id	name	email	avg-rating	- .
---------------	------	-------	------------	-----

Students

student-id	email	phone	grad-year	univ-name
------------	-------	-------	-----------	-----------

classes .

class-id	name	date	
----------	------	------	--

Mentors

mentor-id	name	email	company	
-----------	------	-------	---------	--

mentor-sessions

mentor-session-id	time	duration	mentor-feedback
-------------------	------	----------	-----------------

Student-feedback

Constraints

$\Rightarrow \text{email-id} \Rightarrow \underline{\text{Pk}}$.

↳ Can be a Pk ✓

Reasons for why not to choose email as a Pk

①

Students.

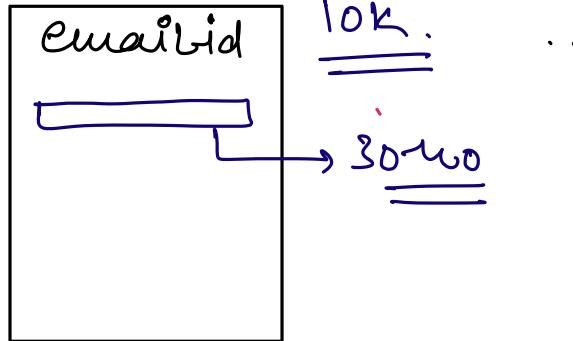
1M

id	email-id	



email-id's

\Rightarrow Space Constraint



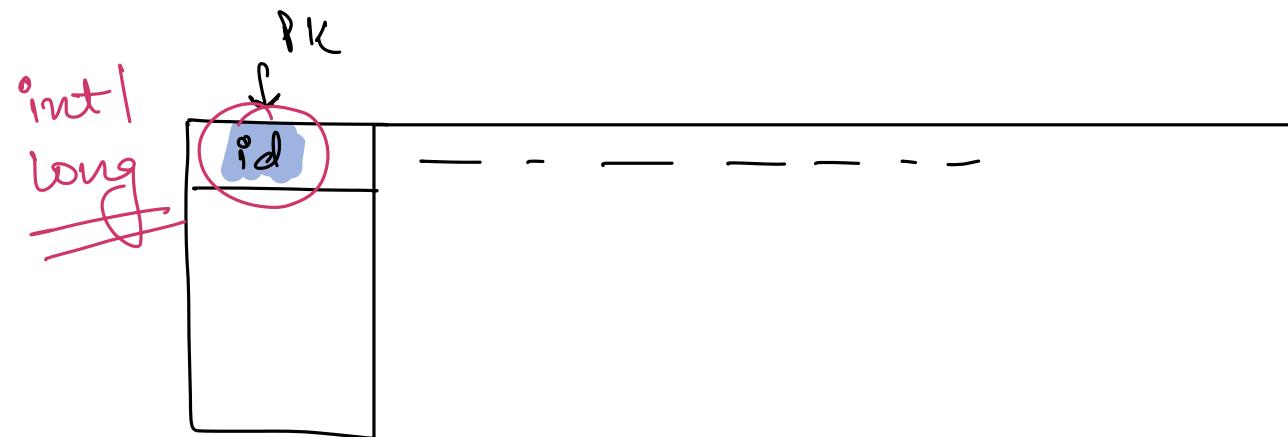
\Rightarrow Rearrangement.

\Rightarrow User's data can change.

(By default DB creates an Index on PK).

Sorted.

\Rightarrow Strings expensive compared int.



\Rightarrow how to represent relations.

CARDINALITY

A B

\Rightarrow how many A's are related to how many B's.

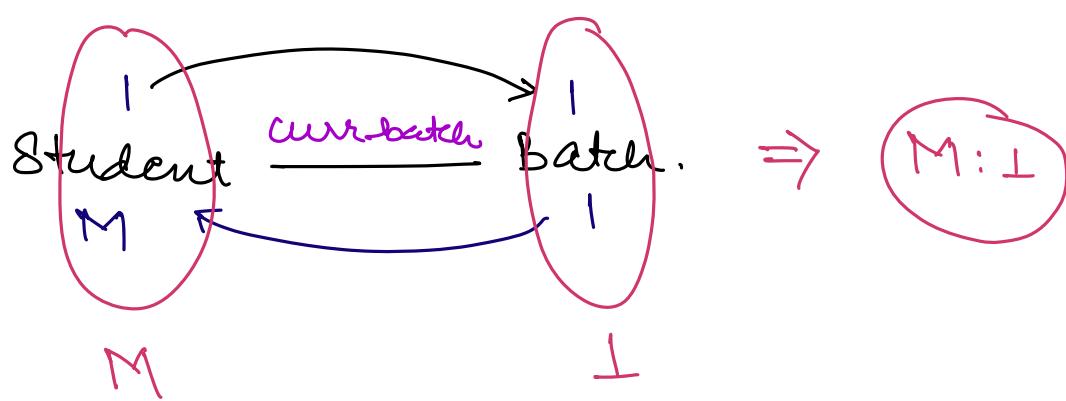
1:1 \Rightarrow 1 (A) is related to 1 (B)

1:M \Rightarrow 1 (A) is related to Many B's

M:1 \Rightarrow 1 (B) is related to Many A's

M:M \Rightarrow Many A's are related to Many B's

eg



\Rightarrow Which relation?

①: An entity can be associated to $[0, \infty]$
of other entity

M: ≥ 1 $[2, \infty)$

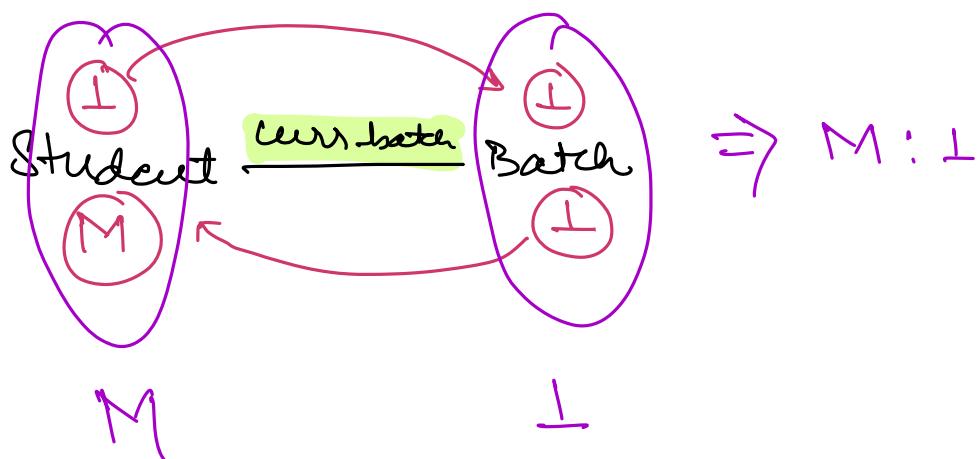
\Rightarrow How to find Cardinality

① Define the relation b/w the entities for
which we want to find cardinality

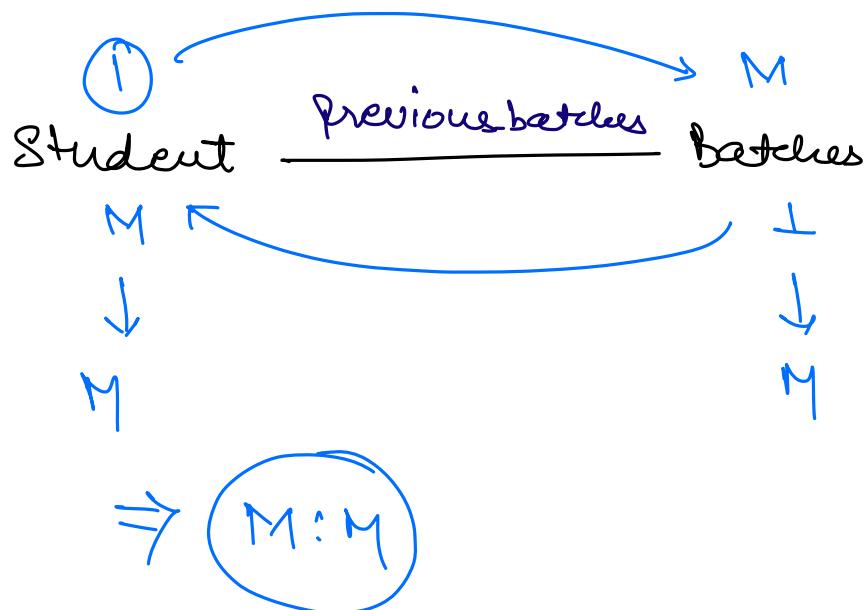
② 2 step approach

i) Write ① on left side & find how many
of right side entities are associated to
entity of left side

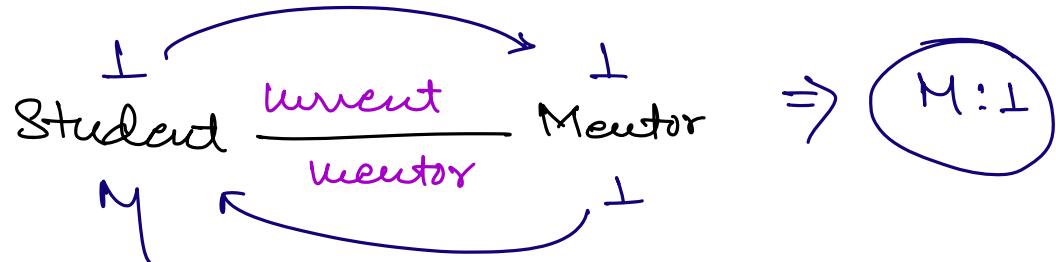
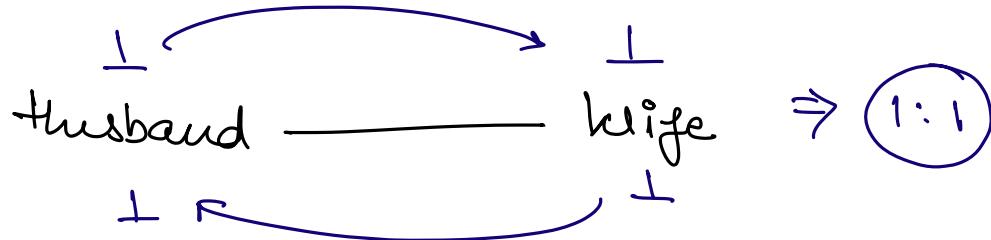
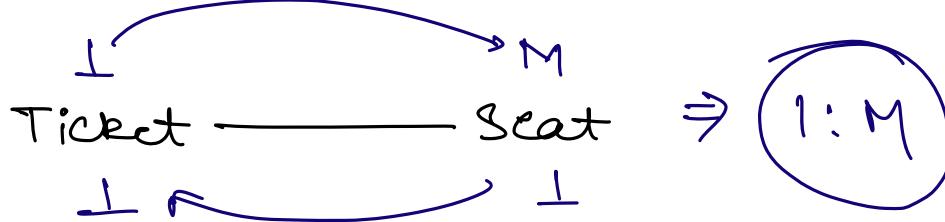
ii) Write ① on right side & find how many
of left side entities are associated to
entity of right side

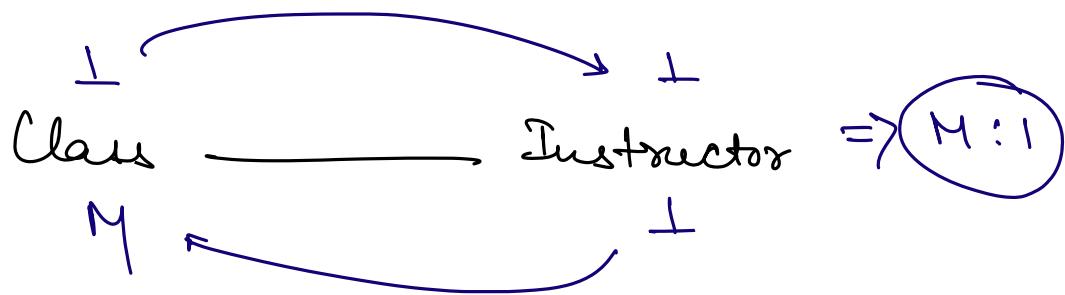


→ If there's any M on a side then take
else take \perp



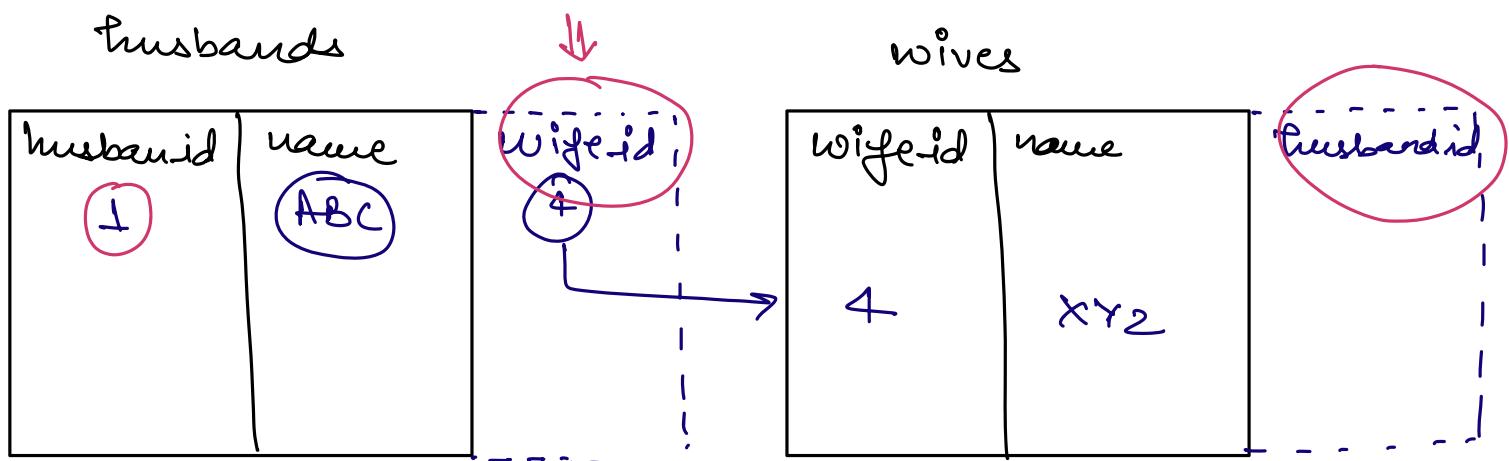
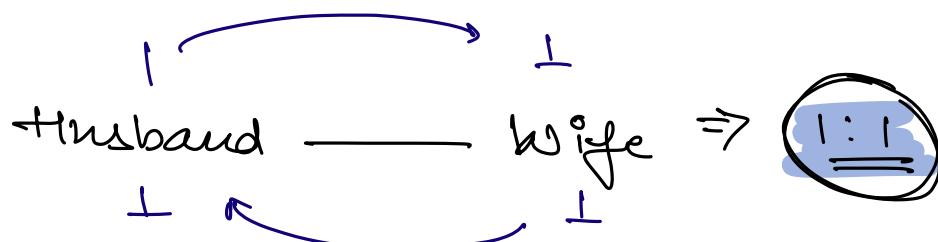
BookMyShow.





\Rightarrow how to represent cardinalities.

1) 1:1



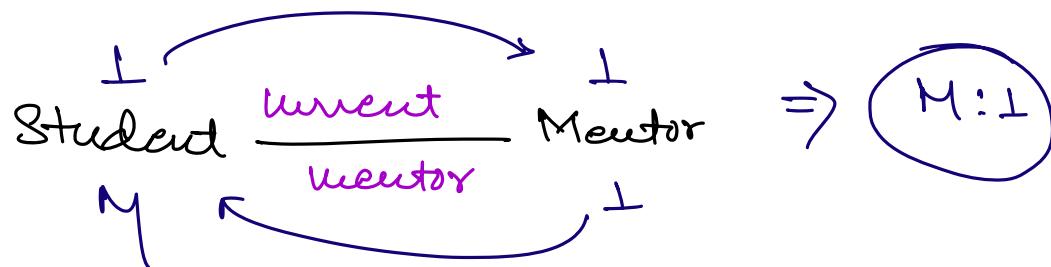
husbandid in wives table.

OR

wife-id is husbands table.

- \Rightarrow Can we do both? Yes
- \Rightarrow Should we do both? \Rightarrow No
- Space.
Redundancy.
- \Rightarrow If we are storing same data at 2 different places then it can lead to inconsistency
- \Rightarrow Update Anomaly

2) $1:M \quad / \quad M:1$



Students

\downarrow ①

id	name	mentor_id
1	Dines	10

mentors

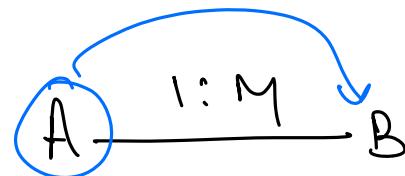
② X X

mentor_id	student_ids
5	[2, 4, 6, 7, 8]

\Rightarrow Every cell should be Atomic.

Student $\xrightarrow{M:1}$ Mentor

\Rightarrow Id of ① side on M side.



\Rightarrow Id of ① in ② table.

3) $M:M$.



Orders

id	Product-ids
1	[10, 11, 12]

Products

id	Order-ids
5	[5, 6, 7..]

\Rightarrow Mapping table | lookup table.

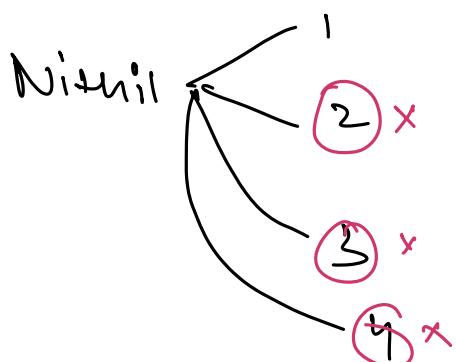
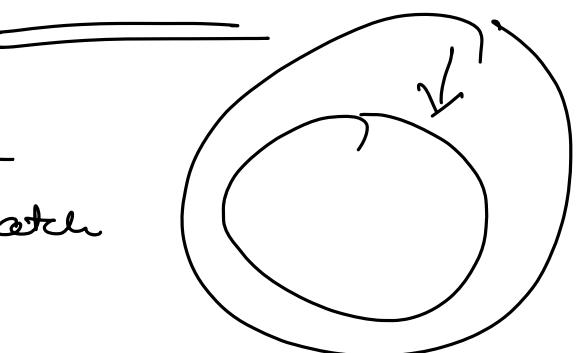
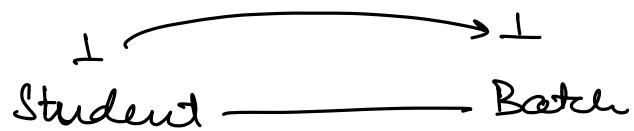
Orders_products.

order_id	product_id
1	10
1	11
1	12
2	10
3	10
4	10

— * —



draw.º.



<u>id:</u> —
Batch: <u>1</u> <u>APR/22</u>
X

* $\vdash M \mid M : \perp$

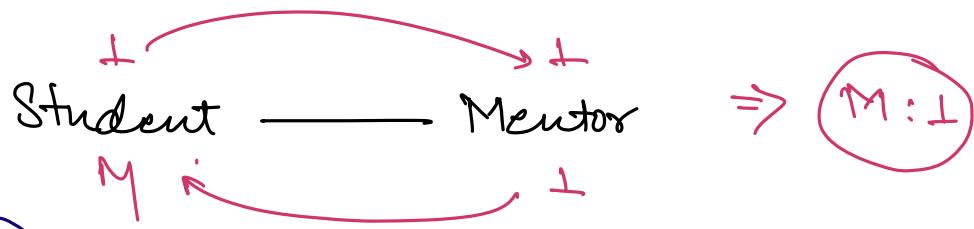


Diagram illustrating the mapping of the many-to-one relationship between Students and Mentors:

Students	$\downarrow \text{id}$	mentors
id	name	mentor_id
1	A	+
2	B	NULL
3	C	5
4	D	2
5	E	NULL
		NULL

Annotations:

- Students table has $10B$ rows.
- Mentors table has $4B$ rows.
- id column has $2M$ distinct values.
- mentor_id column has $30k$ distinct values.
- name column has $2M$ distinct values.

$\Rightarrow \underline{\underline{30k}}$

$\Rightarrow \underline{\underline{19470000}}$

$\Rightarrow 1.947M$ rows will have mentor_id as NULL.

\Rightarrow Sparse relation.

\Rightarrow Space wastage.

$4B + 10B + 4B$ $4B + 10B + 4B$ $4B + 10B + 4B$

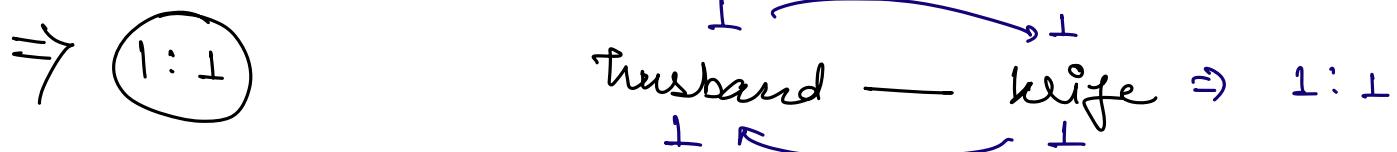
\Rightarrow Mapping / Lookup table.

Students_mentors

Student_id	mentor_id
1	10
2	5
6	4
11	11

\Rightarrow 1:1: No space wastage.

Cores: Joins.



Husband

	wif_id

wife

husband		mid				
id	name	wif-id	marriage-date	eug.-date	status	...
1	A	10				
2	B	4				
3	C	5				
4	D	6				
10						

SRP : Single Responsibility Principle

⇒ Mapping table.

marriage

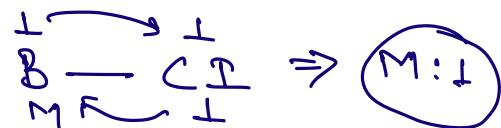
id	h-id	w-id	Marriage date	-	- - -
1	10	5	2021		
		?			

Mapping /
Lookup

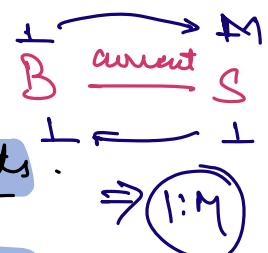
Extra.

Cardinalities	Normal	Sparse	Rel ⁿ has attribute
1 : 1	Id of one side on other side		
1 : M M : 1	Id of one side on <u>M</u> side.	Mapping Table.	Mapping Table.
M : M	Mapping table.		

SCALER.



1. Scaler will have multiple **batches**, for each batch we'll store its name, start month and current instructor.
2. Each batch will have multiple students.
3. Each batch will have multiple **classes**.
4. For each class we need to store their name, date & time, instructor.
5. For each student we'll store their name, email, phone, grad year, university name, etc.
6. Every student will have a buddy, who is another student.
7. A student can be moved from one batch to another batch.
8. For each batch a student goes to we'll store the start date of the batch.
9. Every student can choose their **mentor**, for mentor we'll store their name, email & current company.
10. We'll store the information regarding **mentor sessions** like session-id, time, duration, student, mentor, student rating, mentor rating



11. for every batch we'll store if it is Academy
& DSML

Tables.

batches

batch-id	name	start month	current-instructor-id
----------	------	-------------	-----------------------

instructors

instructor-id	name	email	avg-rating	-
---------------	------	-------	------------	---

Students

student-id	email	phone	grad-year	univ-name	batch-id
------------	-------	-------	-----------	-----------	----------

Classes.

class-id	name	date	instructor-id
----------	------	------	---------------

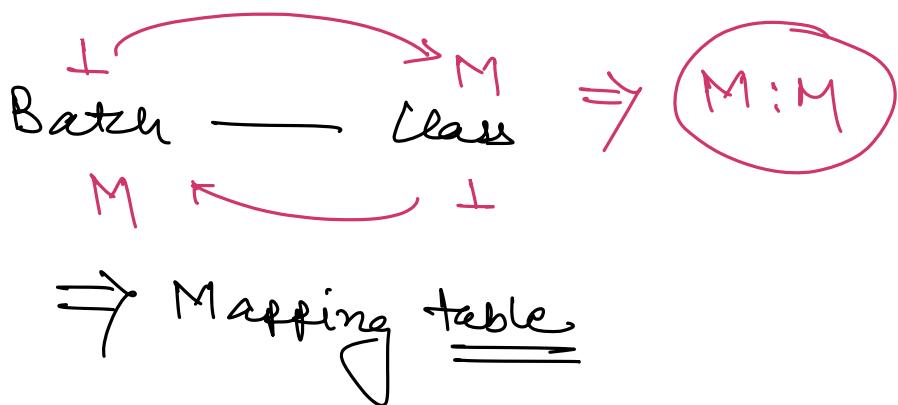
mentors

mentor-id	name	email	company
-----------	------	-------	---------

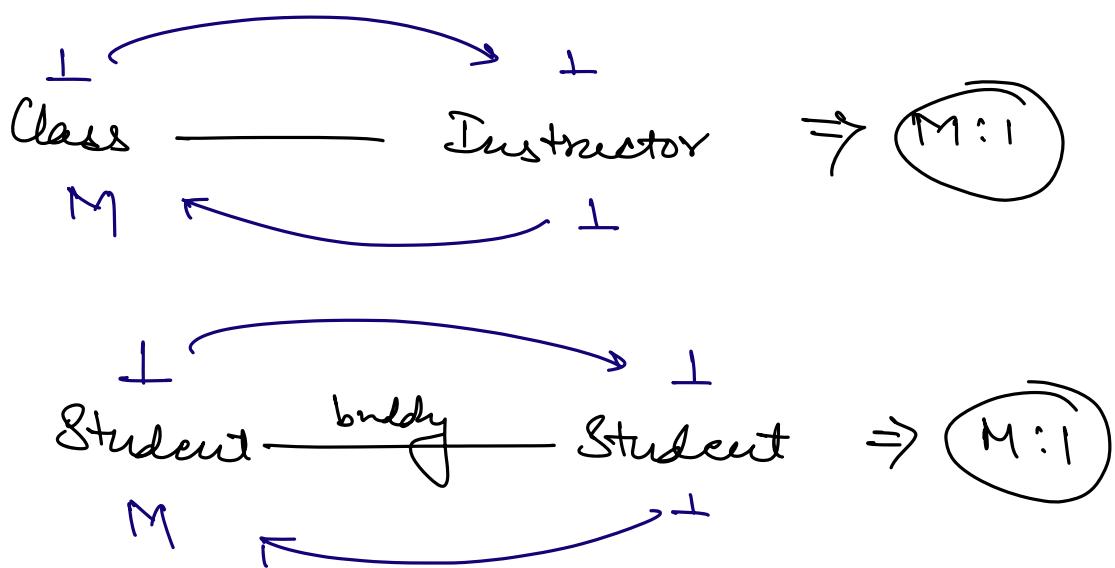
mentor-sessions

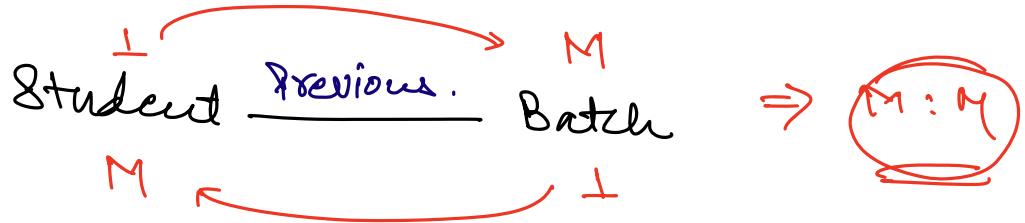
mentor-session-id	time	duration	mentor-feedback
-------------------	------	----------	-----------------

Student-feedback.



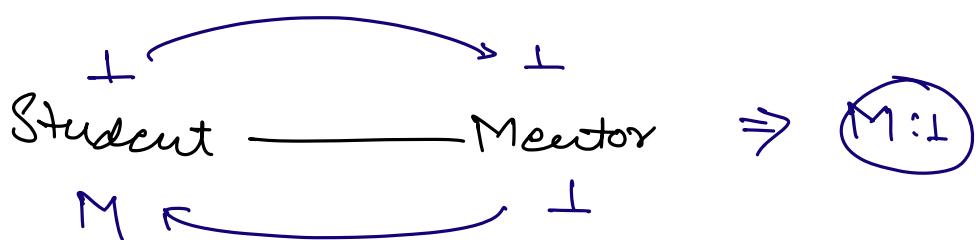
batch_id	Classid





Students_Previoues_batches.

std_id	b_id	Start_date.	batch_level_att
1	10	31/01/2022	80
1	12	31/01/2022	90



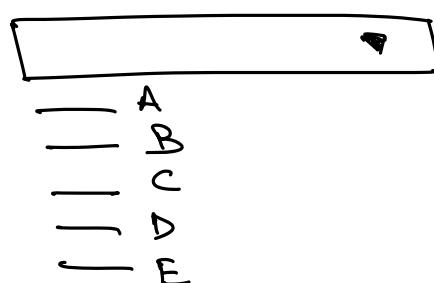
\Rightarrow Drop Down.

Enum

Enum Gender {

MALE, 110

FEMALE 111

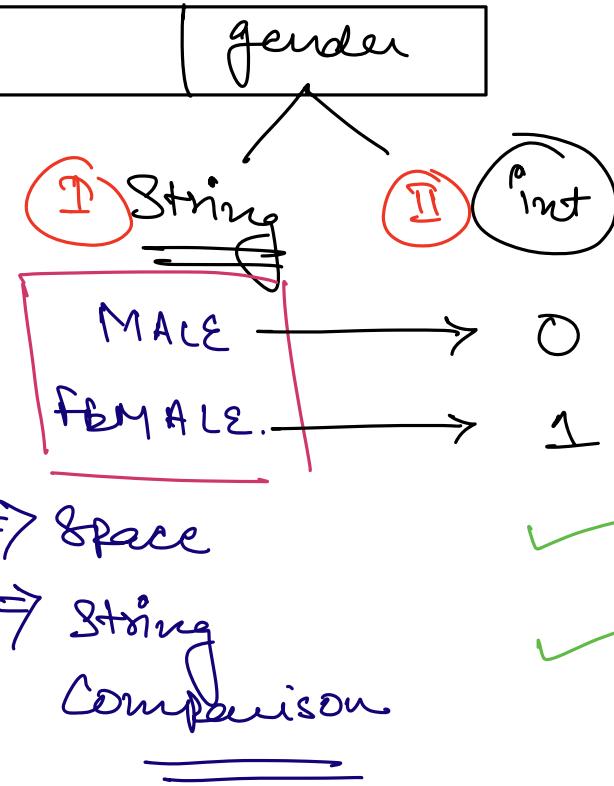


Users.

Enum BatchType {

A, 110
B, 111
C, 112
D, 113
E, 114

3



Enum BatchType {

A 110
B 111
D 112
E 113.

3

ENUM : Predefined set of values

batch-types.

lookup table

id	type
1	A
2	B
3	C
4	D

Students

	batch-type
1	
3	
4	
	<input checked="" type="checkbox"/> NULL

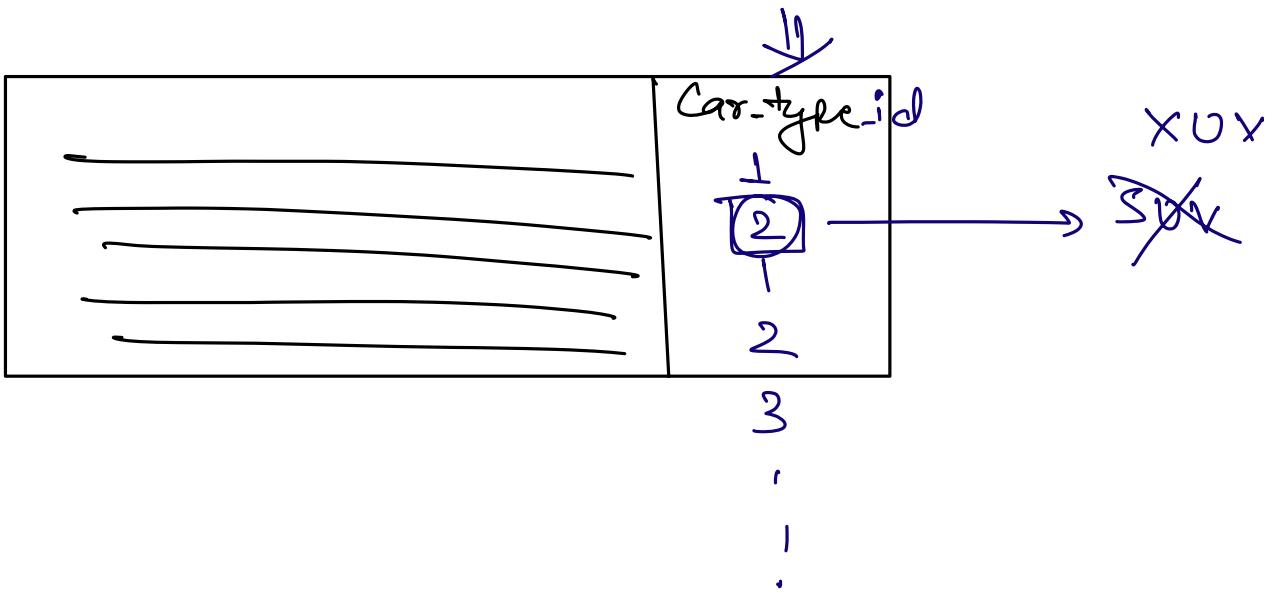
Users.

IM

	Car-type
	SEDAN
	SUV
	HUB
	SUV

Car-type

id	type
1	SEDAN
2	SUV
3	HUB



How to decide PK of Mapping table.

Students_batches.

St_id	b_id	start_date	attendance	...
1	10			
1	20			
1	30			

I (St_id, b_id)

II Create a new id column

<u>id</u>	st-id	b-id	start-date	attendance	<u>pop</u>
1					
2					
3					
4					
5					
:					

⇒ Indexing.

⇒ Index on 2 columns will be bigger than index on 1 column

⇒ If we are taking Composite key as the primary key then it sorts the data based on it, the query to fetch data will be faster.

⇒ Get all the students with pop > 80 in their previous batches.

```
Select *
from _____
where st-id = 10
```

⇒ Composite key

PK
Index.

^o id	St_id	b_id	...
1	20	10	
2	4	5	
3	20	11	
4	4	6	
5	1	2	
6			

Attribute that uniquely identifies a row.

(A, B, C)

PK
Create index.

^o id	St_id	b_id